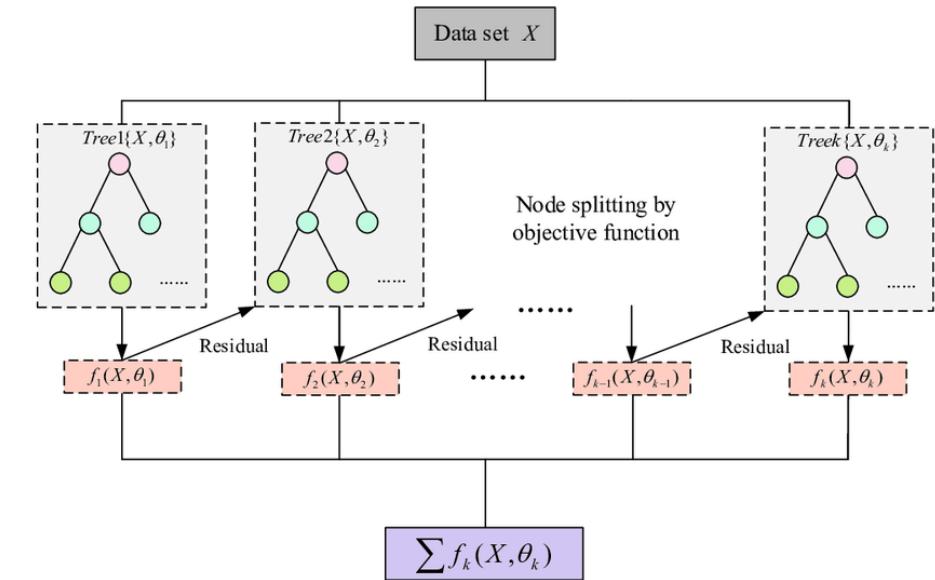
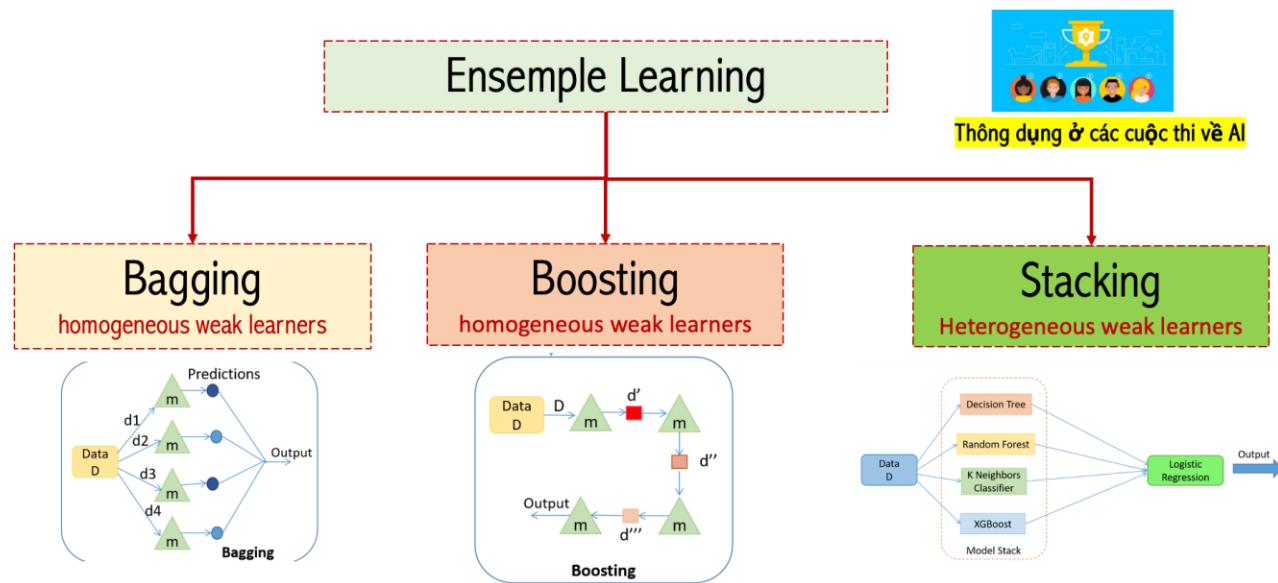


AdaBoost & Gradient Boost

(Basic, Advanced Concepts and Its Applications)



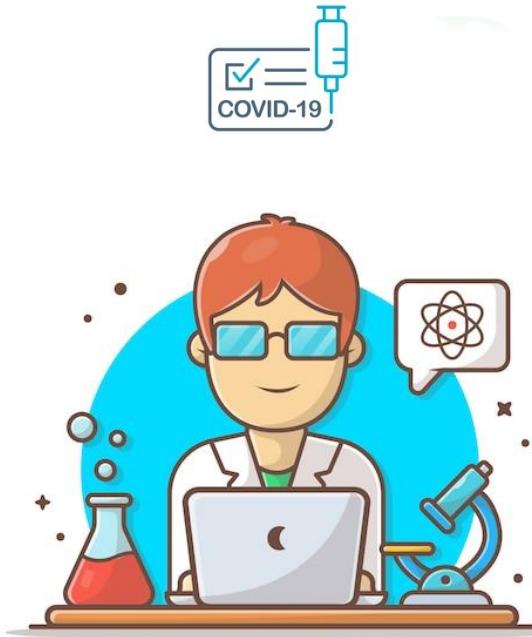
Vinh Dinh Nguyen
PhD in Computer Science

Outline



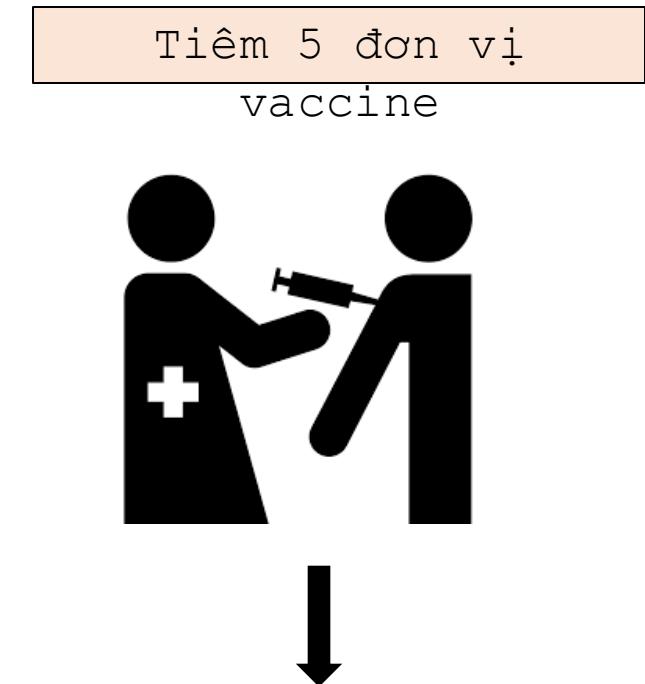
- Boosting Techniques
- AdaBoost Clearly Explain
- Gradient Boost Clearly Explain
- Time Series Data: Predicting Energy Consumption
- Summary

Decision Tree and Its Variance



Unit(đơn vị)	Effect (hiệu quả) (%)
10	98
20	0
35	100
5	44
...	...

Khi có 1 vaccine ra đời, chúng ta muốn dự đoán xem nó hiệu quả bao nhiêu % ứng với từng liều lượng dùng trên bệnh nhân.

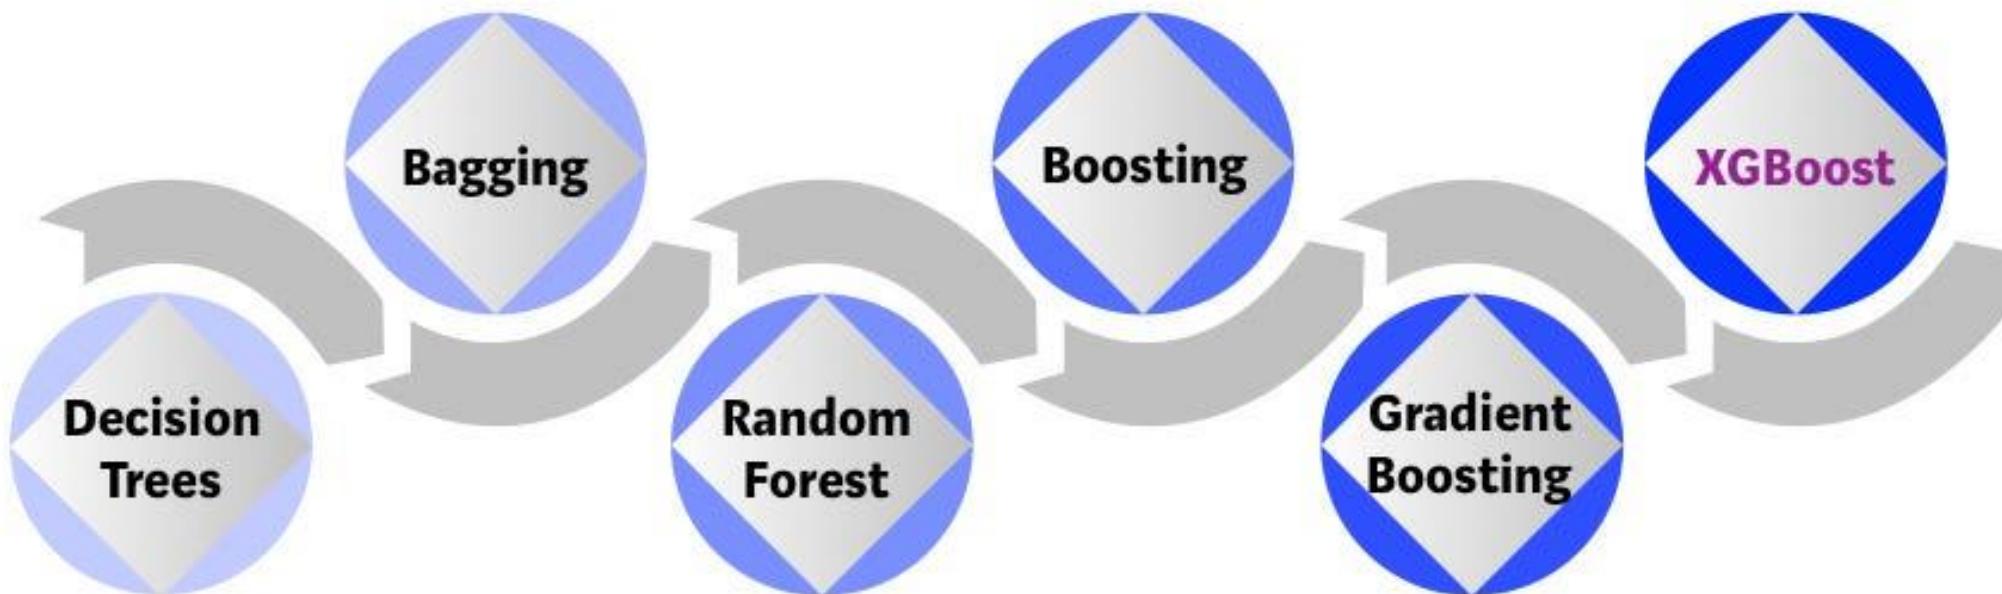


Decision Tree and Its Variance

Bootstrap aggregating or Bagging is a ensemble meta-algorithm combining predictions from multiple decision trees through a majority voting mechanism

Models are built sequentially by minimizing the errors from previous models while increasing (or boosting) influence of high-performing models

Optimized Gradient Boosting algorithm through parallel processing, tree-pruning, handling missing values and regularization to avoid overfitting/bias

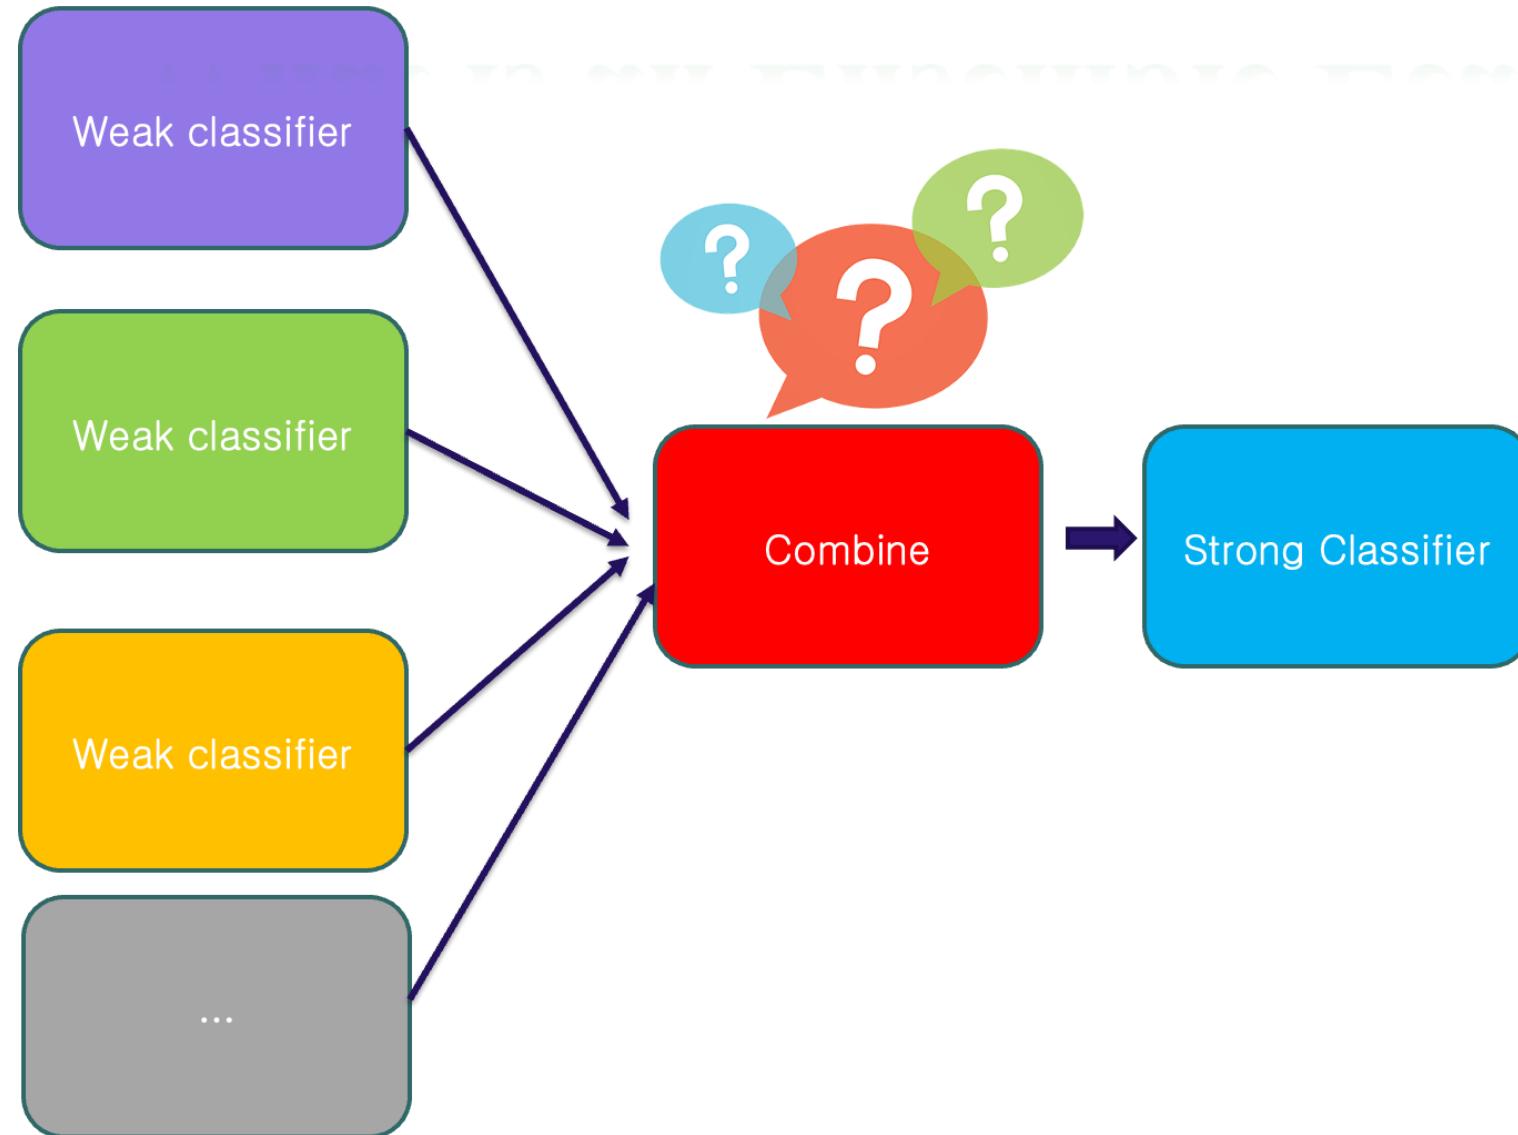


A graphical representation of possible solutions to a decision based on certain conditions

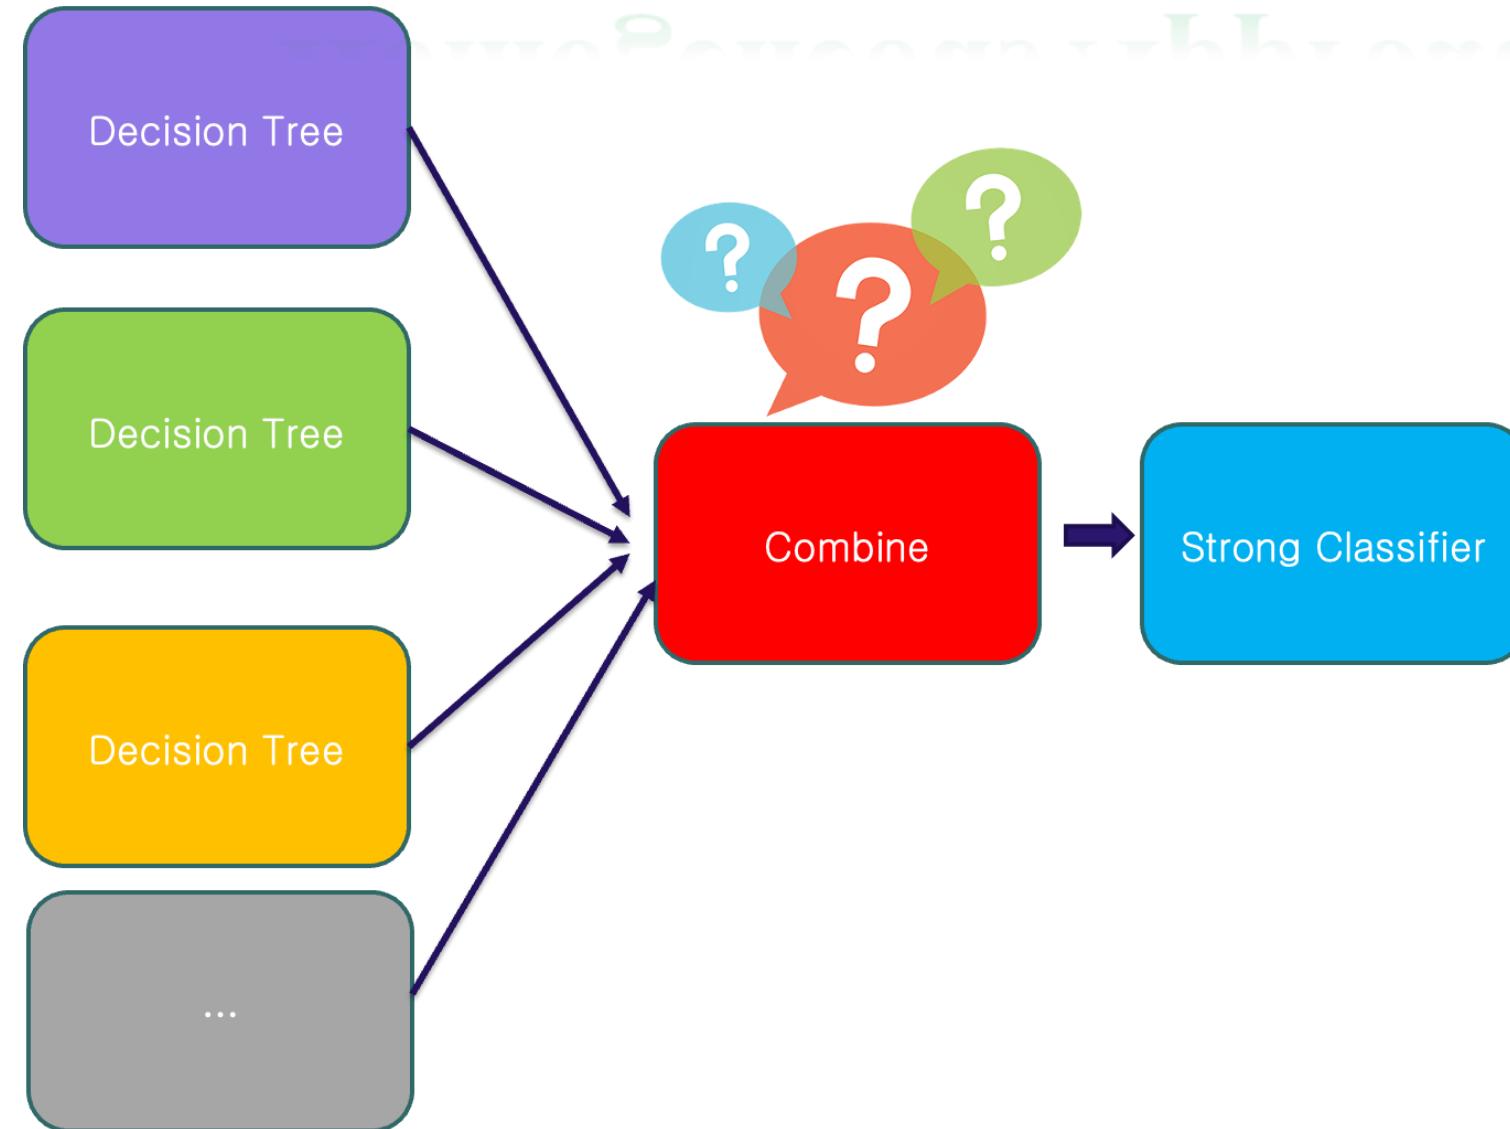
Bagging-based algorithm where only a subset of features are selected at random to build a forest or collection of decision trees

Gradient Boosting employs gradient descent algorithm to minimize errors in sequential models

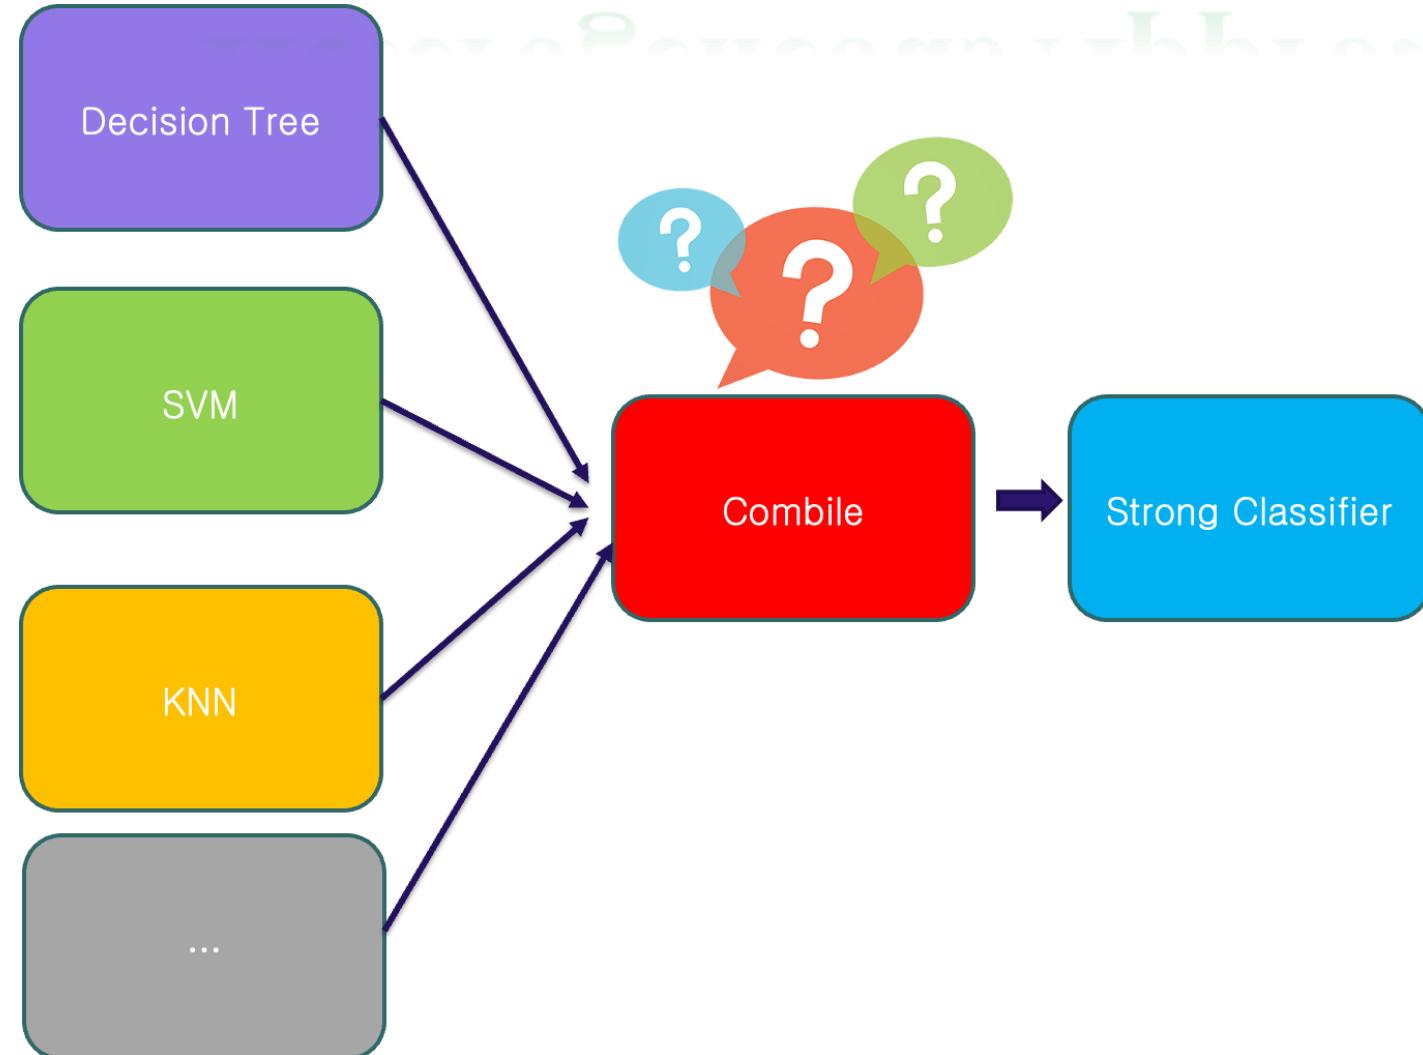
What is an Ensemble Learning?



Homogeneous Approach



Heterogeneous Approach

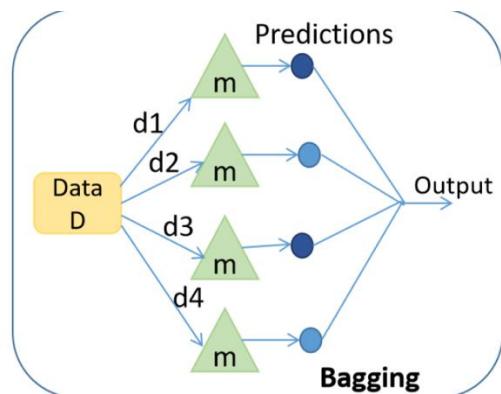


Ensemble Learning Techniques

Ensemble Learning

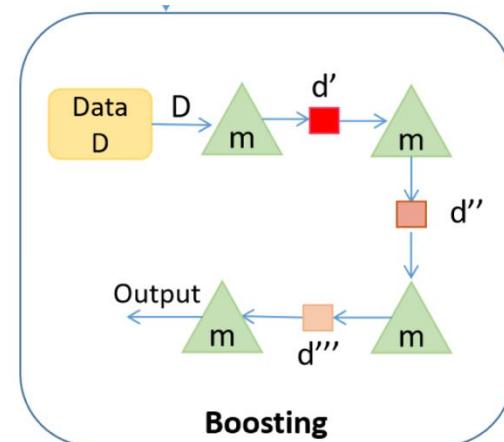
Bagging

homogeneous weak learners



Boosting

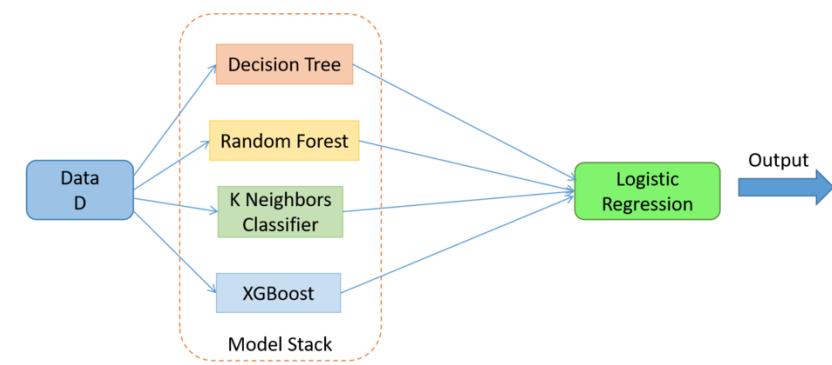
homogeneous weak learners



Thông dụng ở các cuộc thi về
AI

Stacking

Heterogeneous weak learners

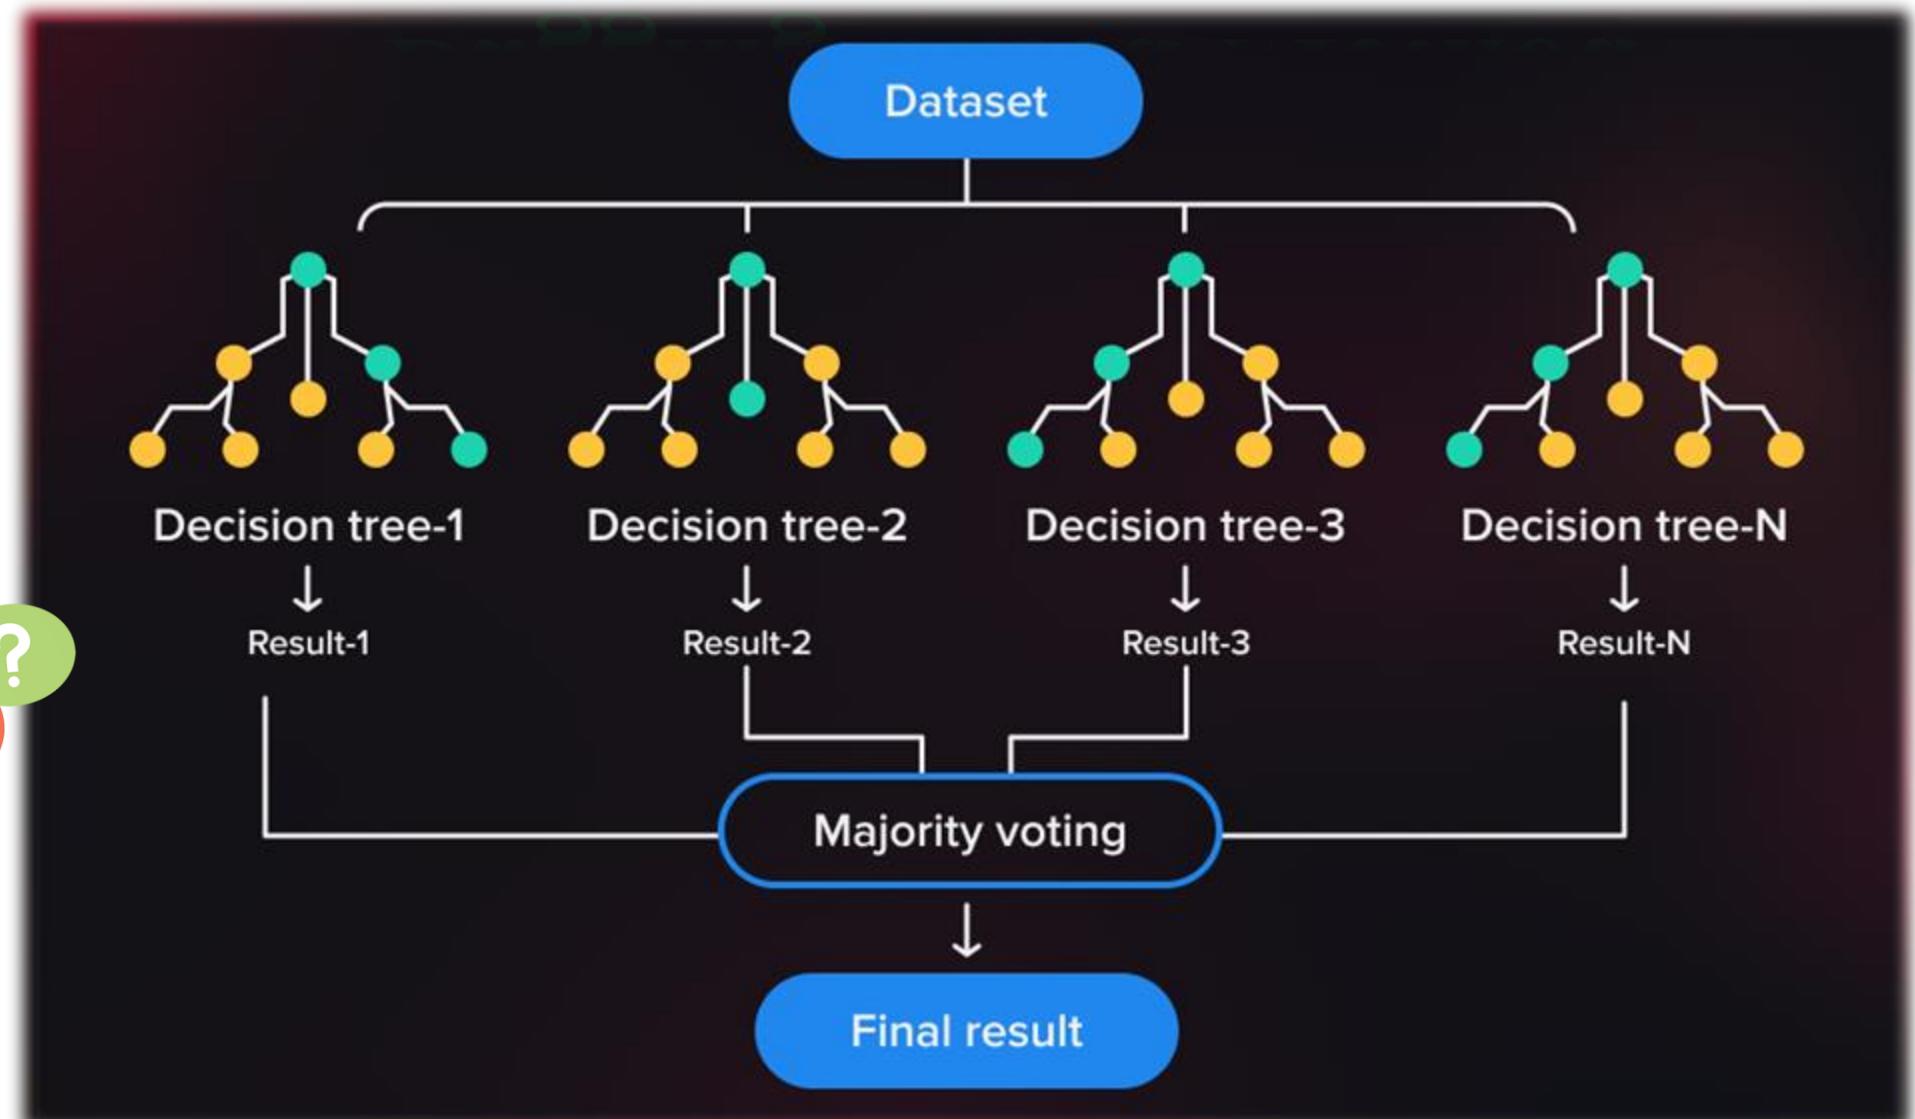


Bagging-based Method

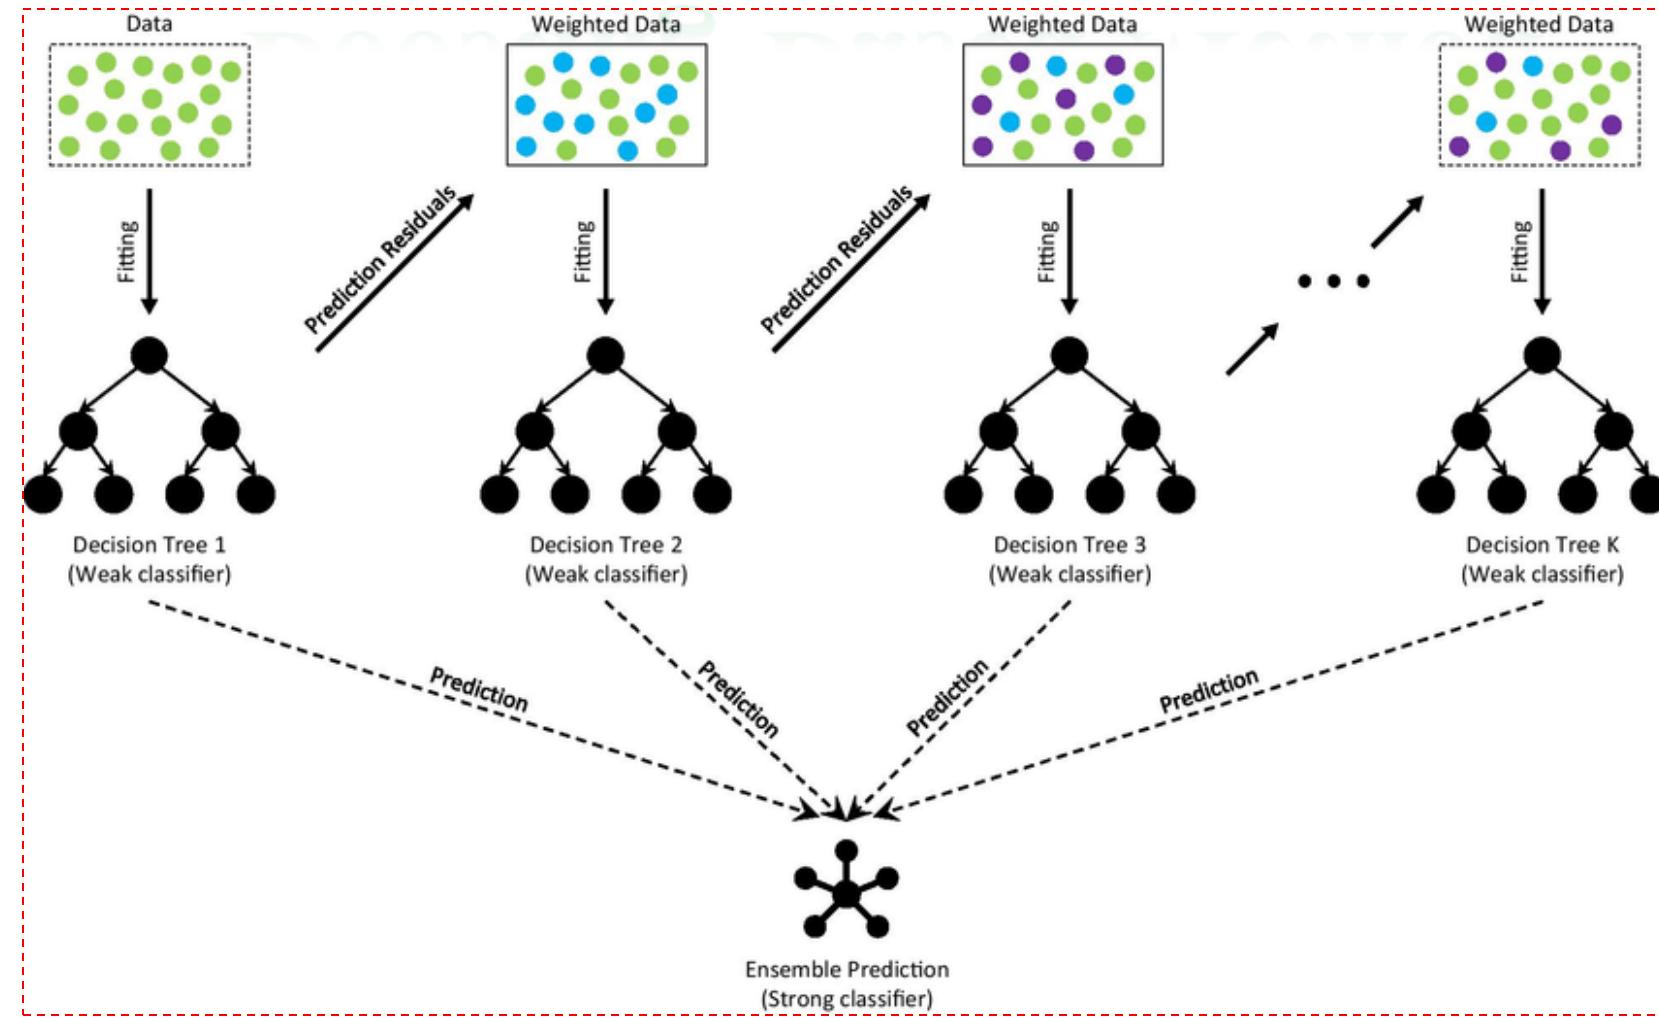
Random
Forest

Last Week

Limitation

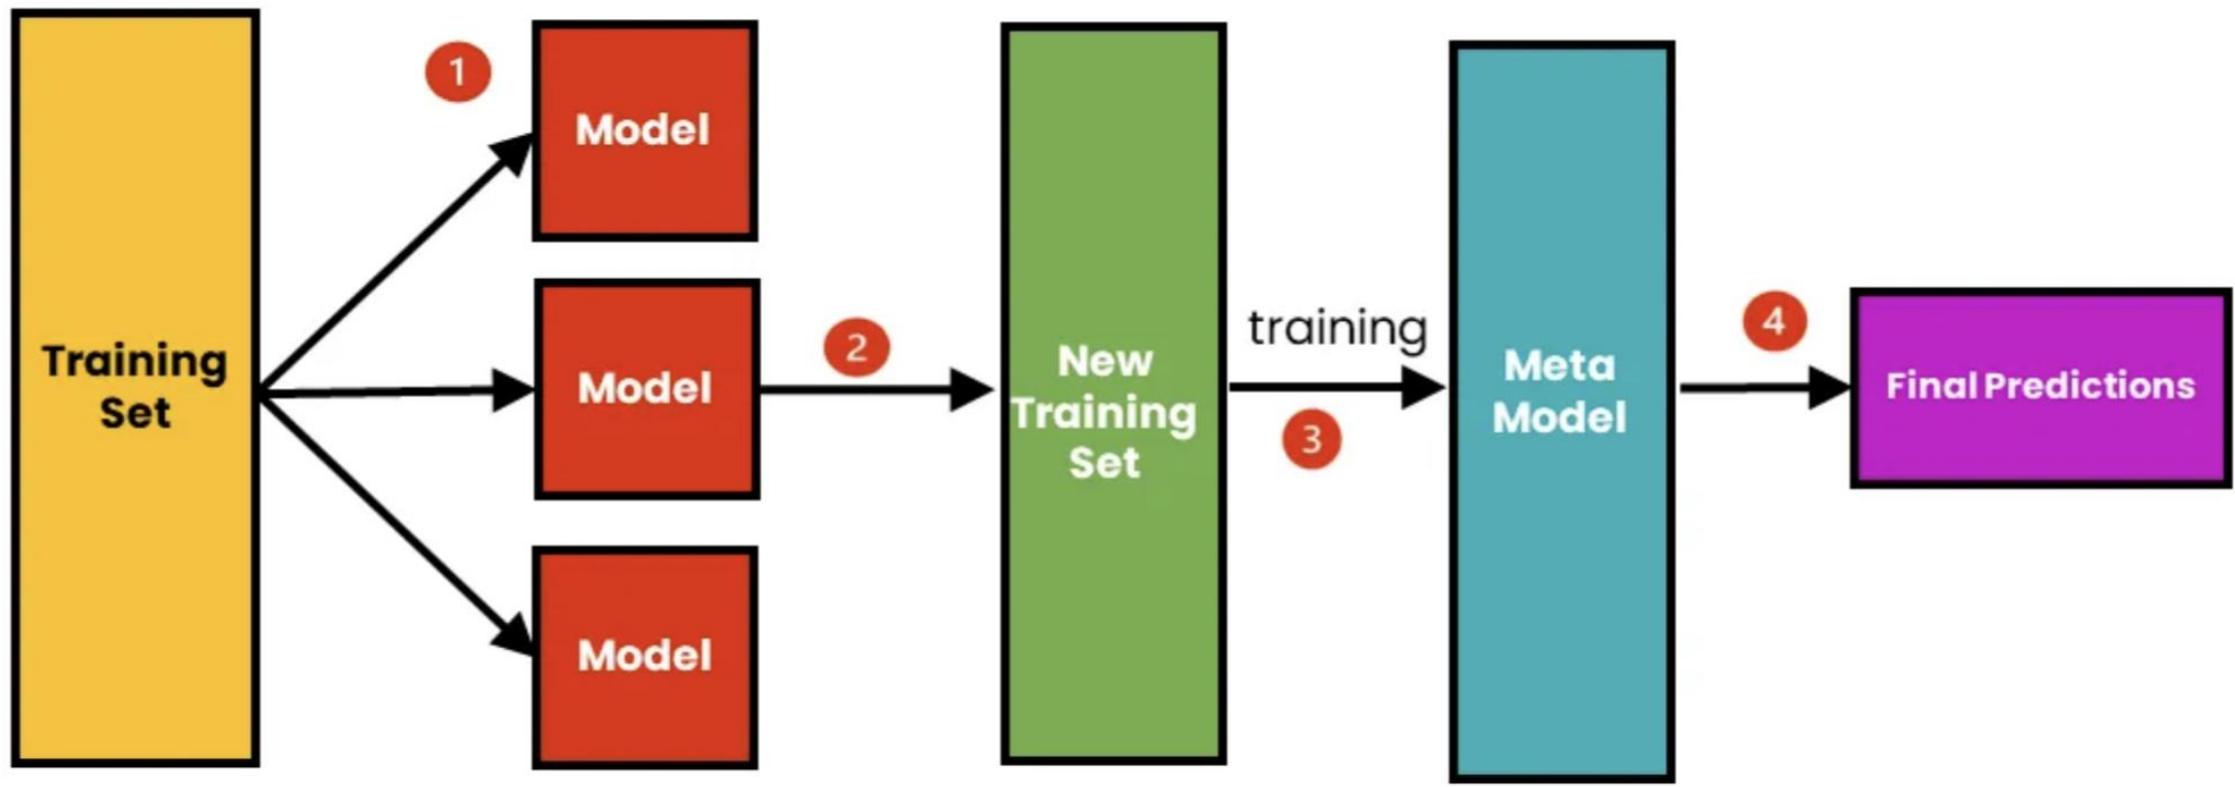


Boosting-Based Method

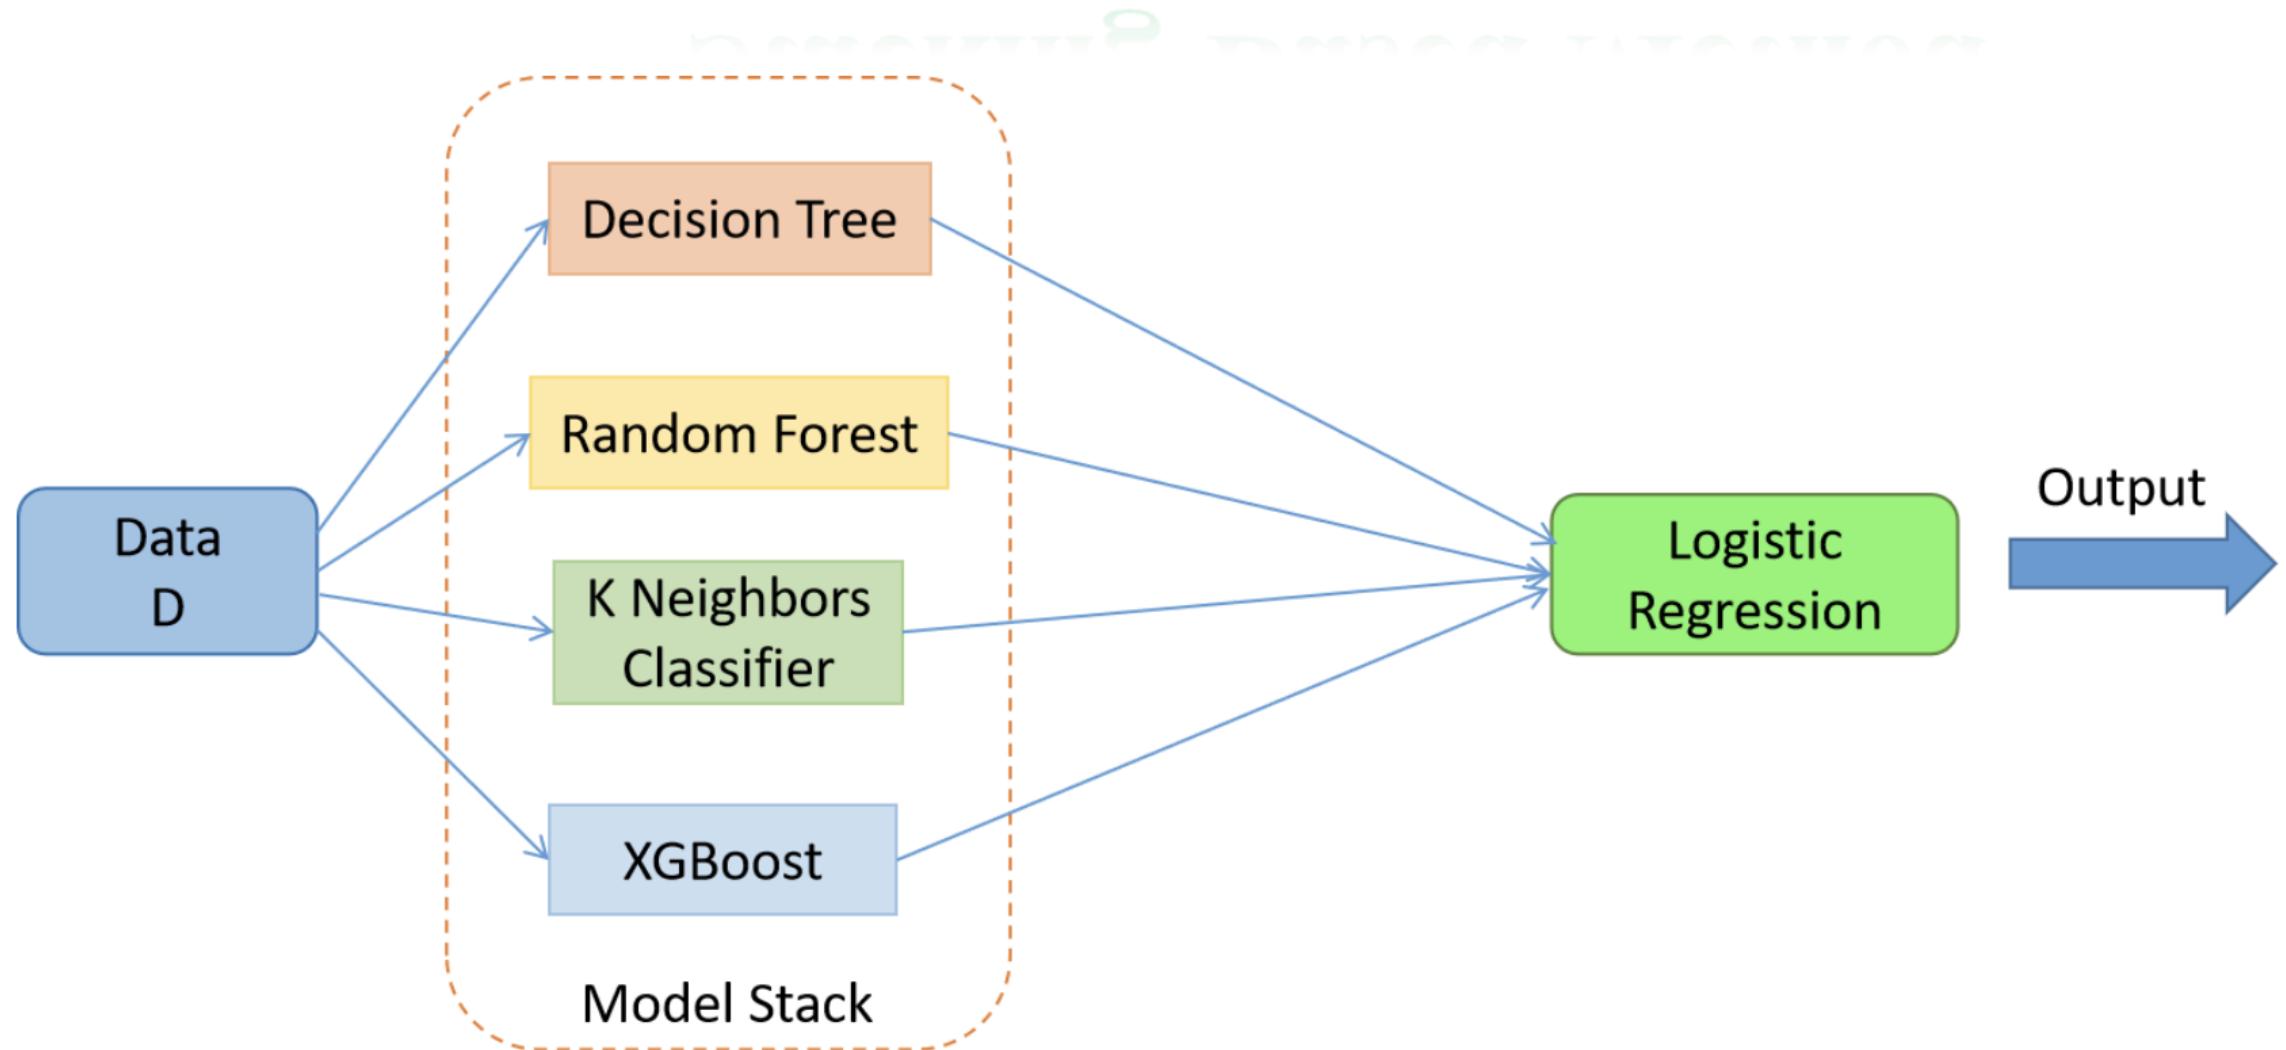


Stacking-Based Method

The Process of Stacking

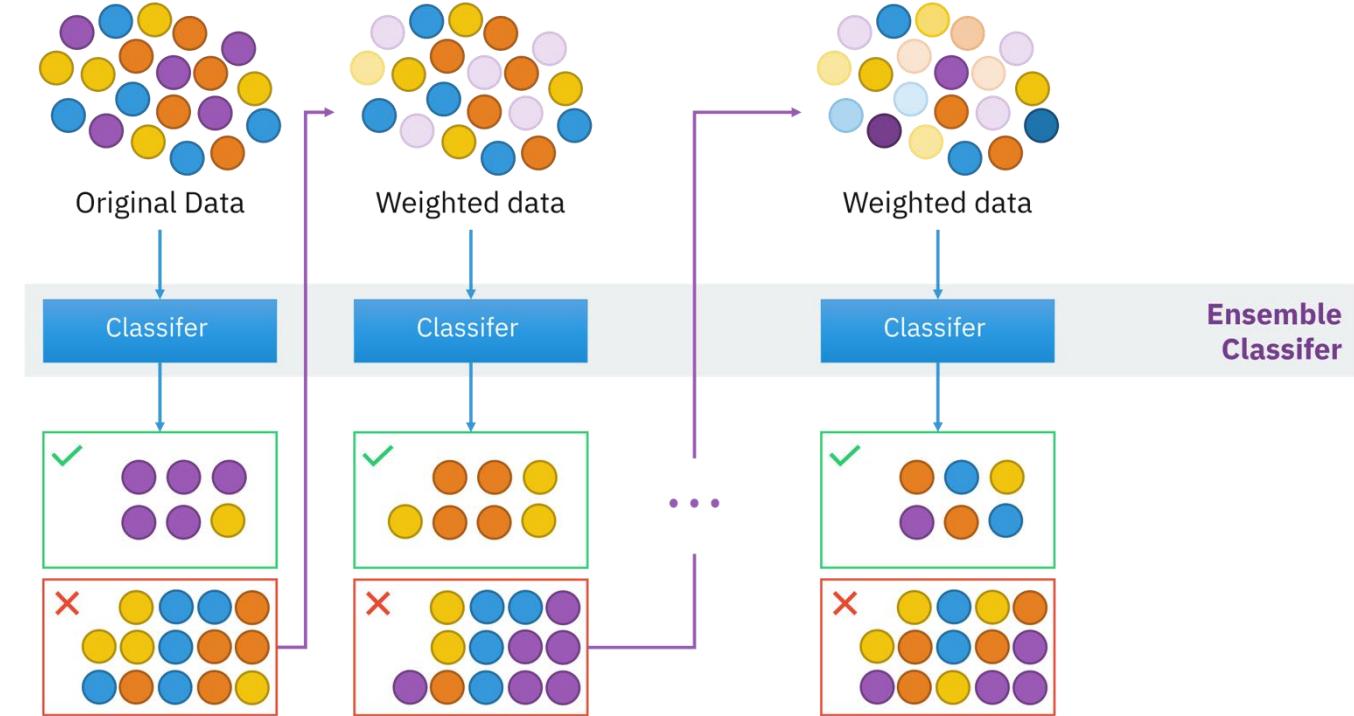


Stacking-Based Method



Boosting Technique

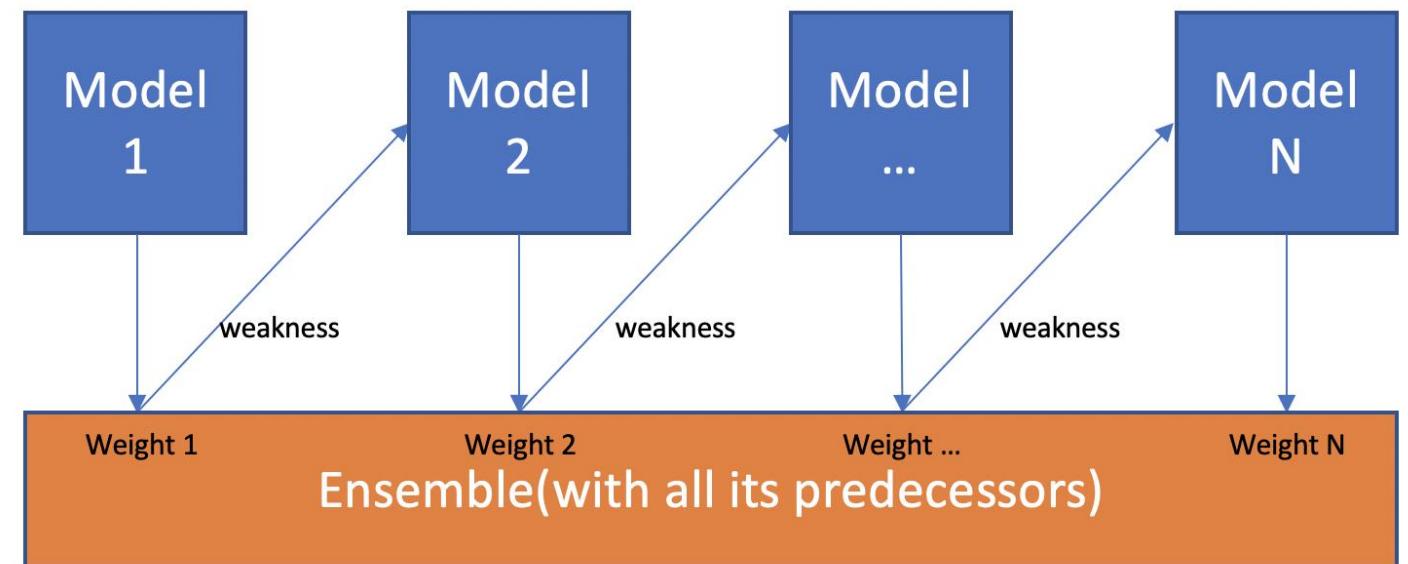
Boosting is an ensemble modelling, technique that attempts to build a strong classifier from the number of weak classifiers



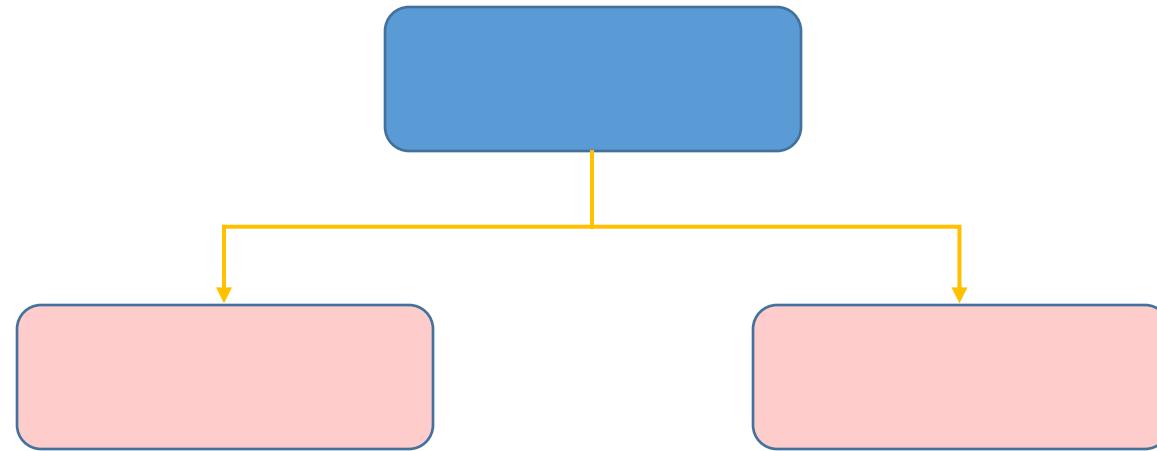
Boosting Technique

Boosting is an ensemble modelling, technique that attempts to build a strong classifier from the number of weak classifiers

Model 1,2,..., N are individual models (e.g. decision tree)



Stump Definition



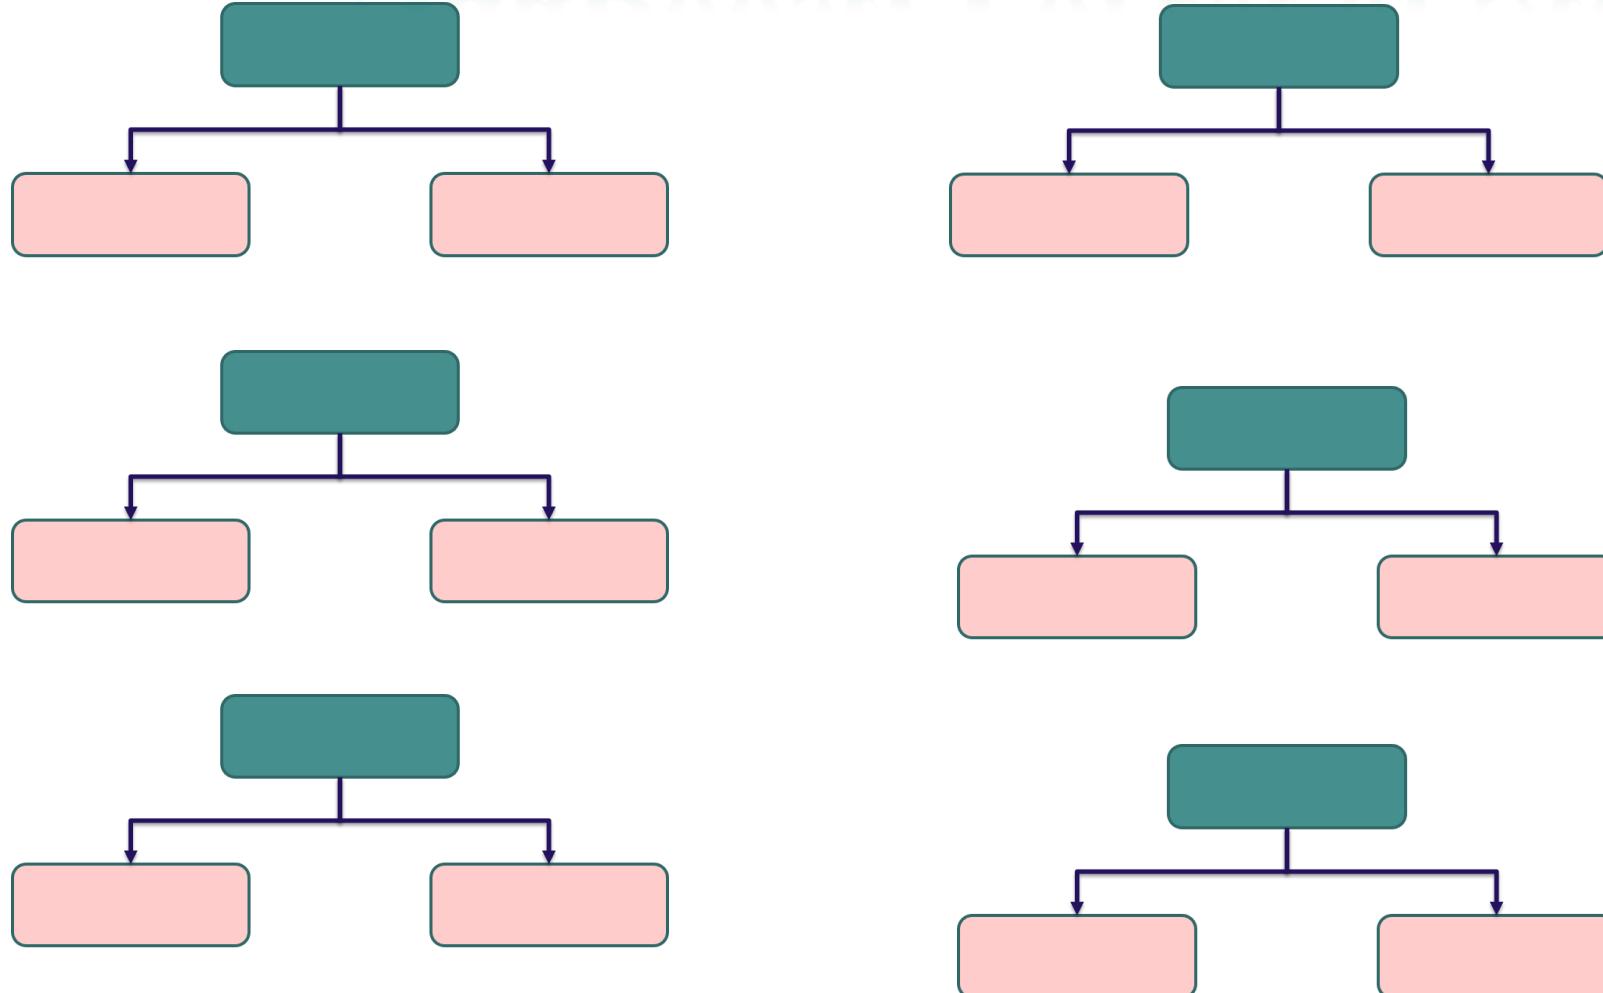
a node with two leaves and this is known as
Stump

Outline

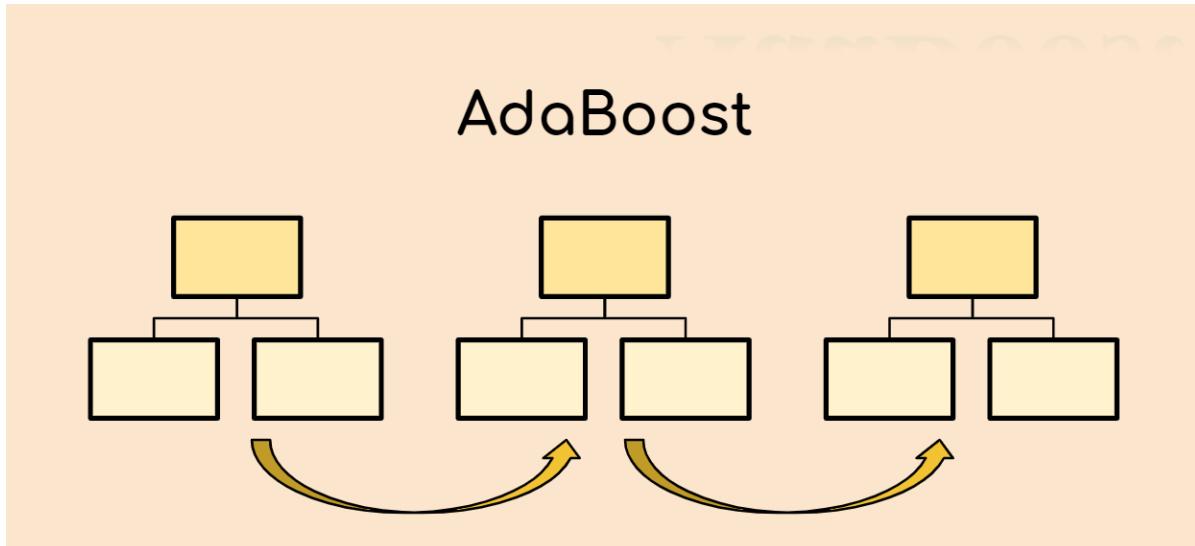


- Boosting Techniques
- AdaBoost Clearly Explain
- Gradient Boost Clearly Explain
- Time Series Data: Predicting Energy Consumption
- Summary

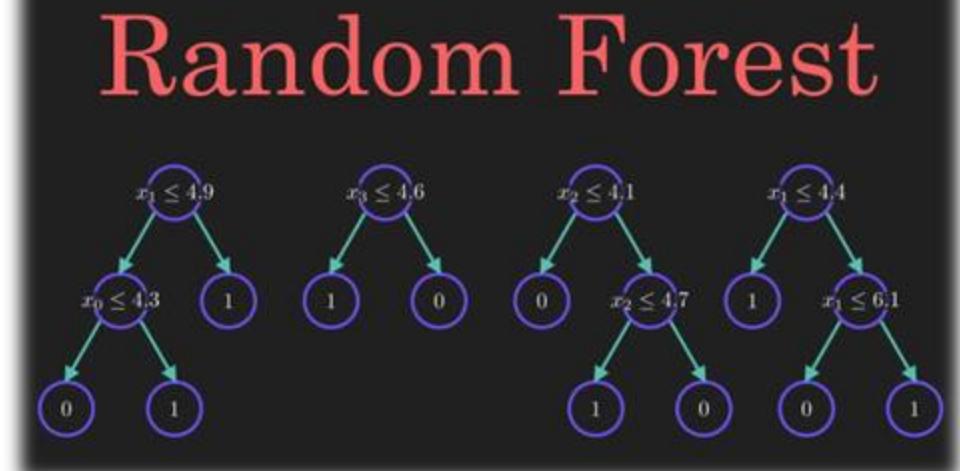
AdaBoost: Forest of Stump



AdaBoost: Forest of Stump



VS

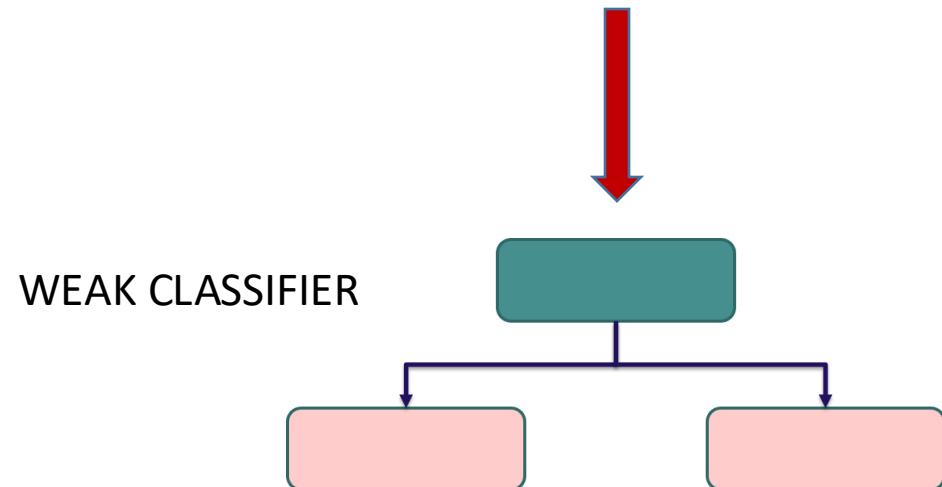
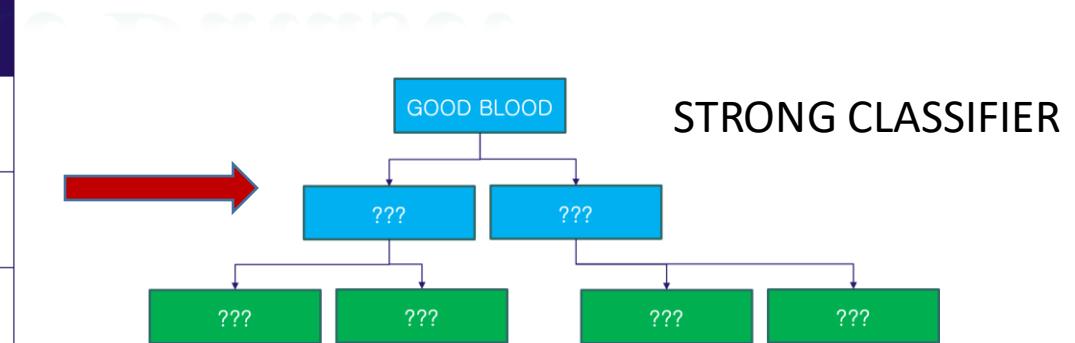


Sample Dataset

CHEST PAIN	GOOD BLOOD CIRCULATION	BLOCKED ARTERIES	WEIGHT	HEART DISEASE
NO	NO	NO	125	NO
YES	YES	YES	180	YES
YES	YES	NO	210	NO
YES	NO	YES	167	YES

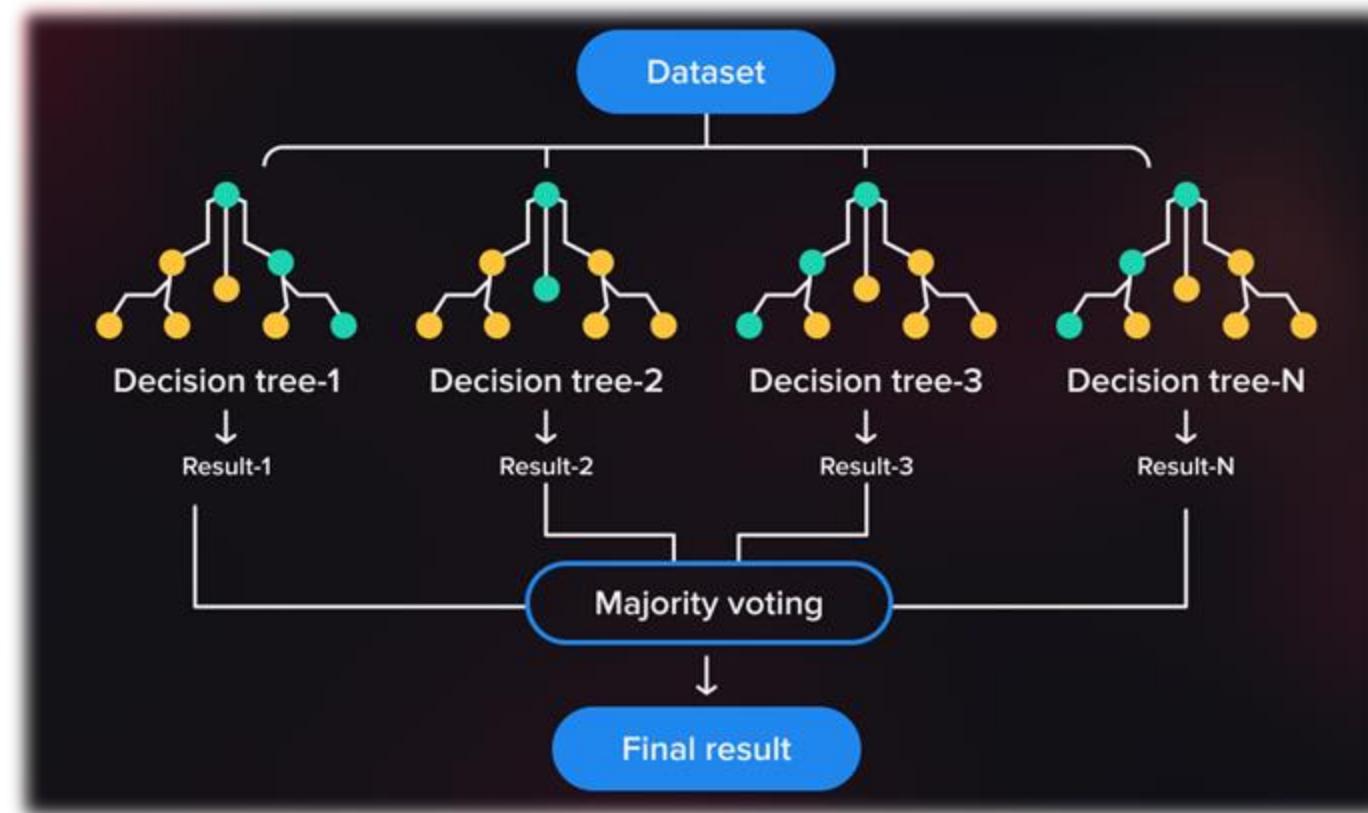
Sample Dataset

CHEST PAIN	GOOD BLOOD CIRCULATION	BLOCKED ARTERIES	WEIGHT	HEART DISEASE
NO	NO	NO	125	NO
YES	YES	YES	180	YES
YES	YES	NO	210	NO
YES	NO	YES	167	YES



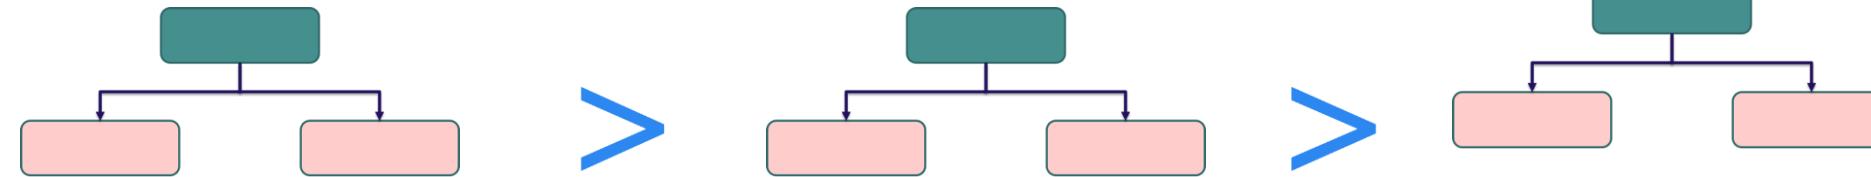
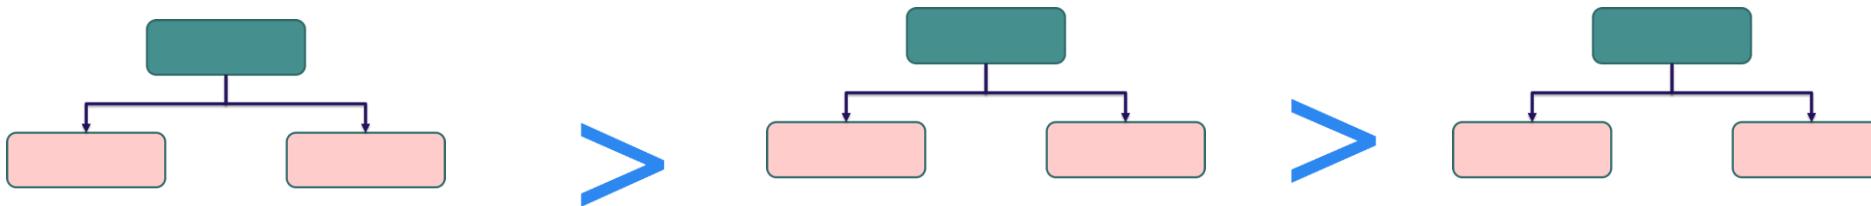
Random Forest

- Each tree in the random forest has equal votes(weights) on the final decision



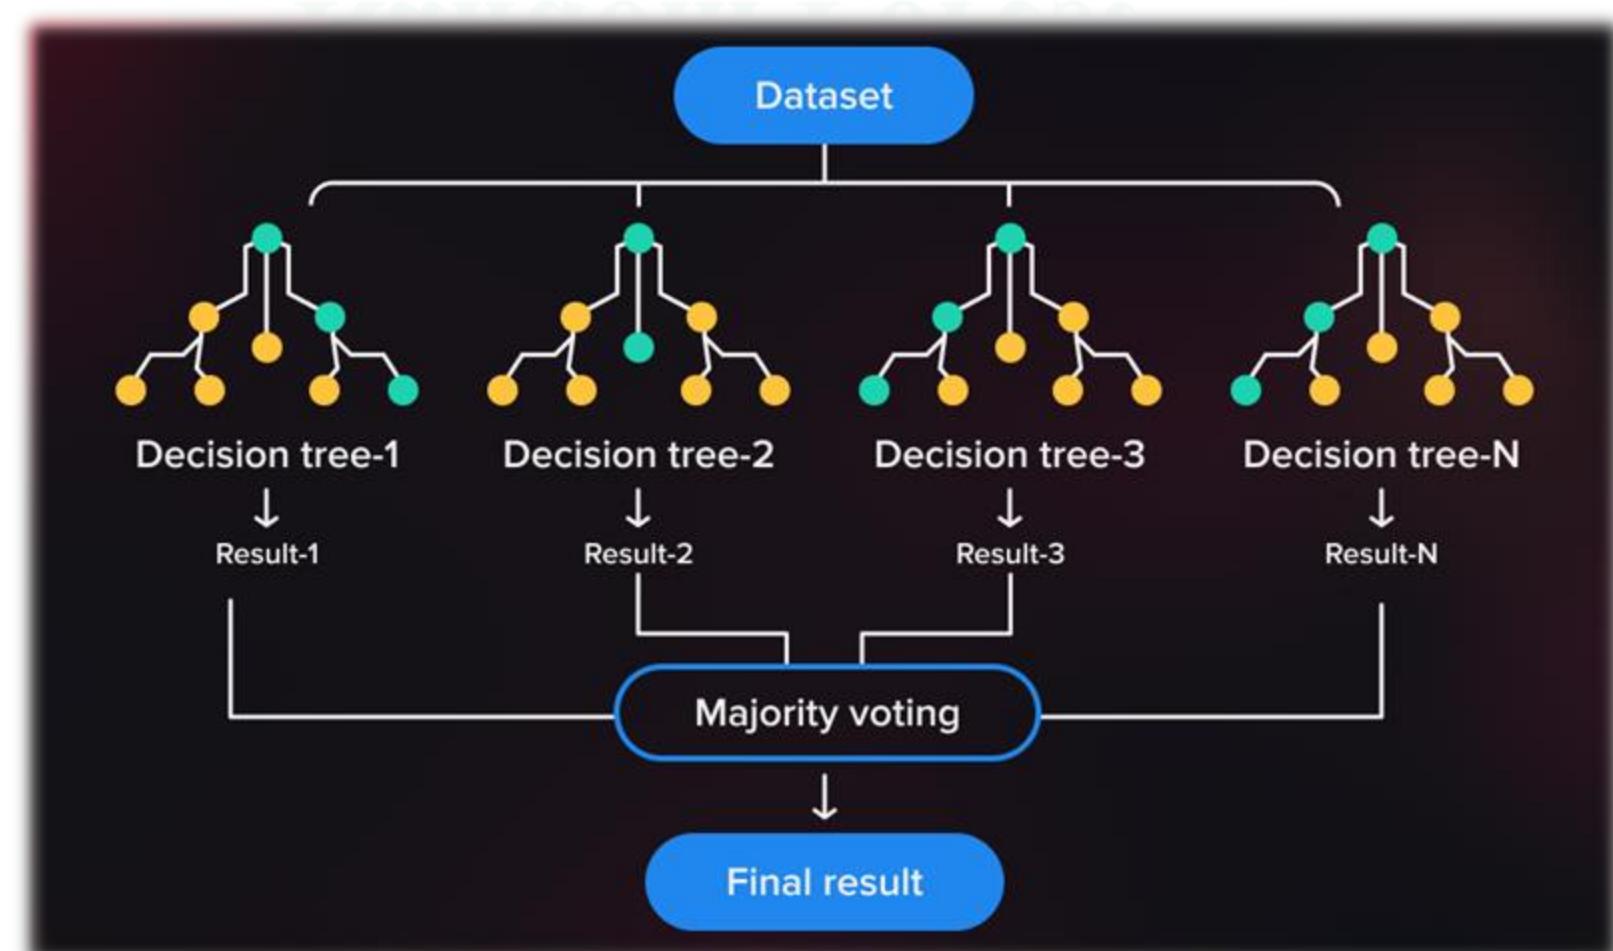
AdaBoost: FOREST OF STUMP

- Stump are not equally weighted in the final decision.
- Stump that create more error will have less contribution in the final decision

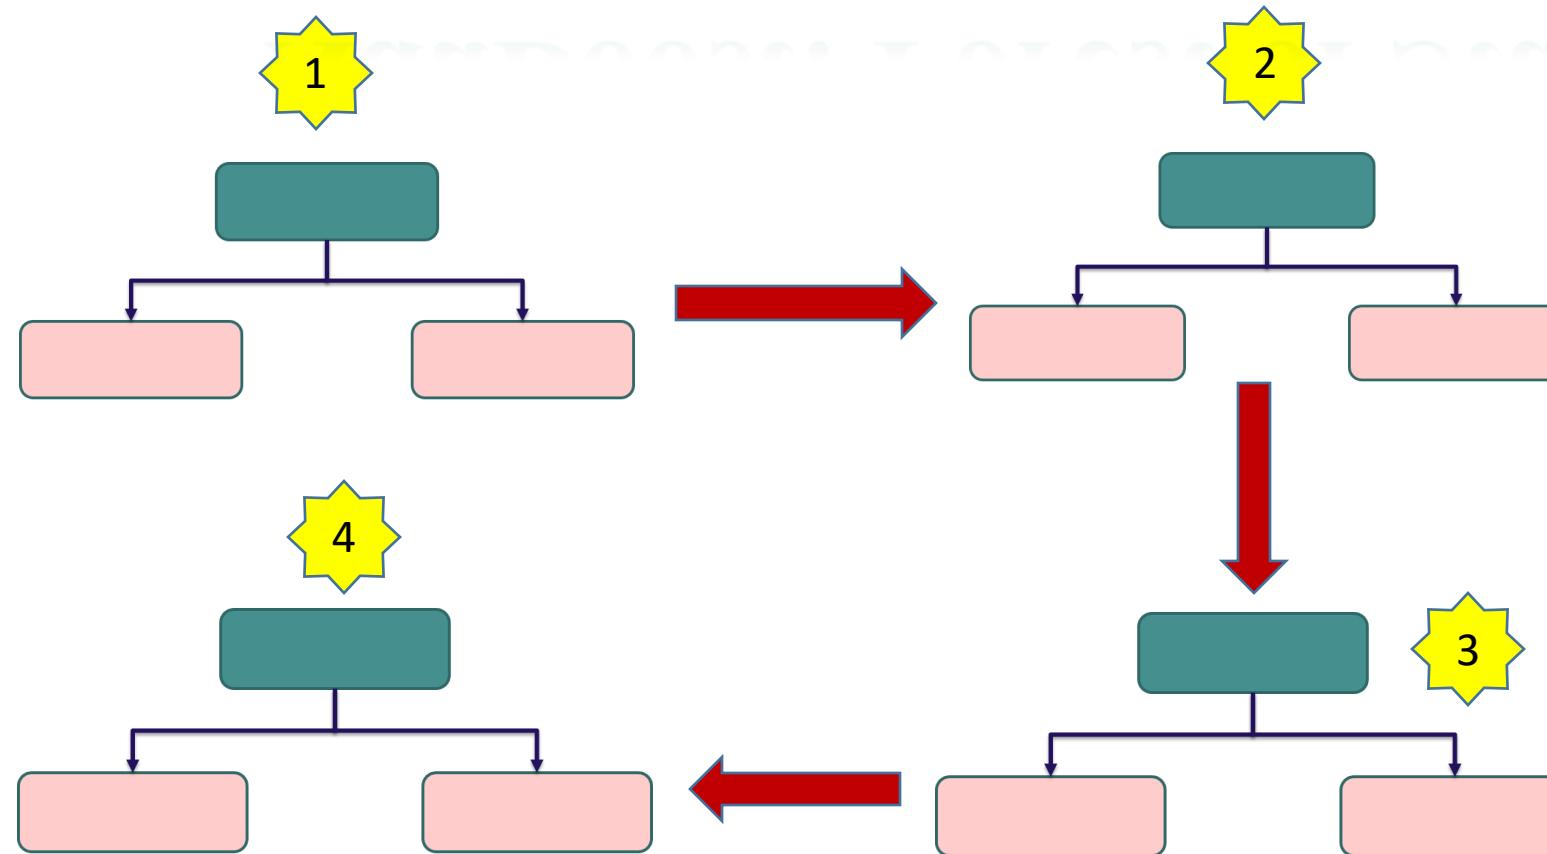


Random Forest

Tree are indepently created



AdaBoost: Forest of Stump



Differences Between RF and AdaBoost

1
ONE

Weak Learners is a ___
AdaBoost combines a lot of ___

2

Stumps have various ___ to
the final result

3

Each stump is created by
considering ___

Heart Disease Dataset

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease
Yes	Yes	205	Yes
No	Yes	180	Yes
Yes	No	210	Yes
Yes	Yes	167	Yes
No	Yes	156	No
No	Yes	125	No
Yes	No	168	No
Yes	Yes	172	No

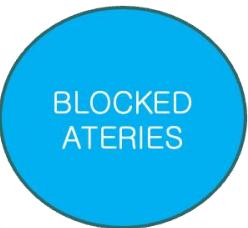
Important of sample = Sample weight = 1 / number of samples = 1/8

1st Stump in the Tree

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease
Yes	Yes	205	Yes
No	Yes	180	Yes
Yes	No	210	Yes
Yes	Yes	167	Yes
No	Yes	156	No
No	Yes	125	No
Yes	No	168	No
Yes	Yes	172	No



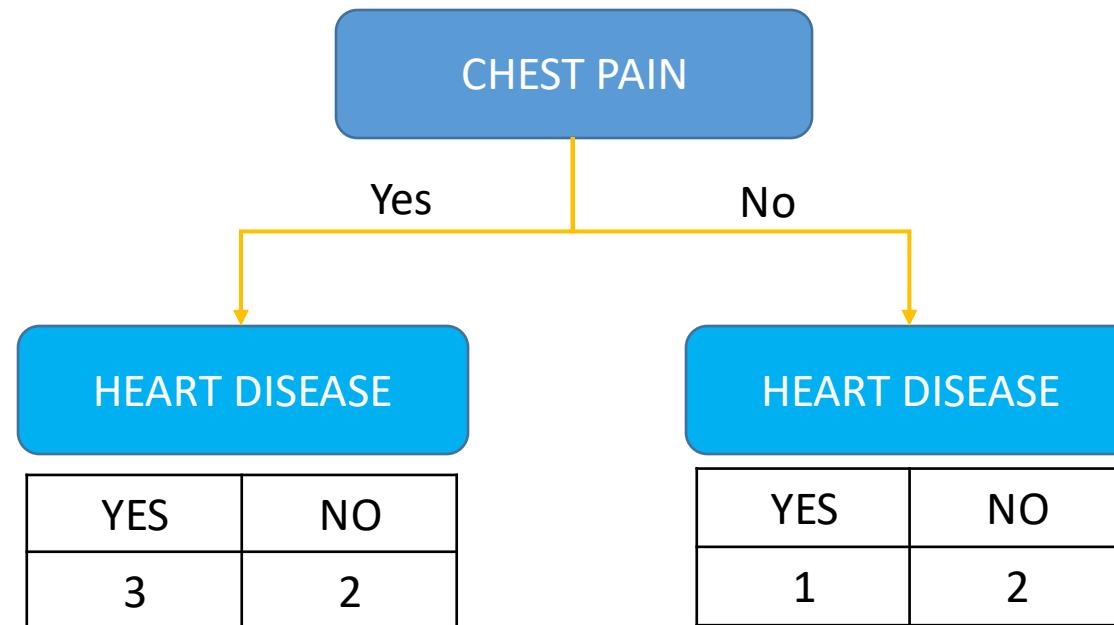
WHICH
ONE?



Compute Gini Index For Chest Pain

Chest Pain	Heart Disease	Sample Weight
Yes	Yes	1/8
No	Yes	1/8
Yes	Yes	1/8
Yes	Yes	1/8
No	No	1/8
No	No	1/8
Yes	No	1/8
Yes	No	1/8

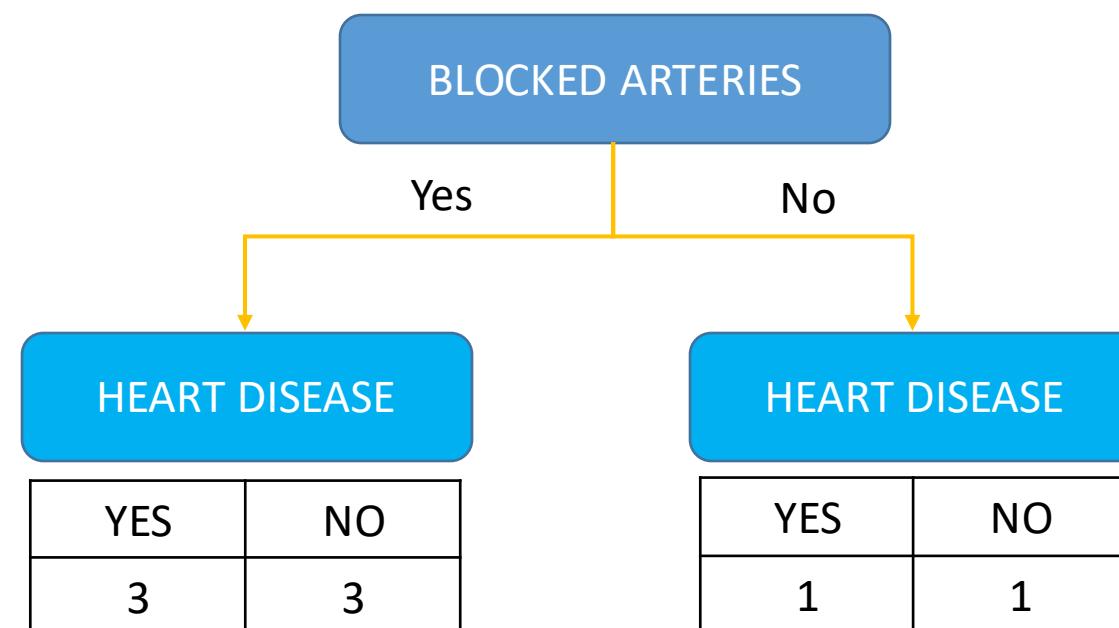
Gini index = $5/8 * (1 - (3/5)^2 - (2/5)^2) + 3/8 * (1 - (1/3)^2 - (2/3)^2) = 0.57$



Gini Index for Blocked Arteries

Blocked Arteries	Heart Disease	Sample Weight
Yes	Yes	1/8
Yes	Yes	1/8
No	Yes	1/8
Yes	Yes	1/8
Yes	No	1/8
Yes	No	1/8
No	No	1/8
Yes	No	1/8

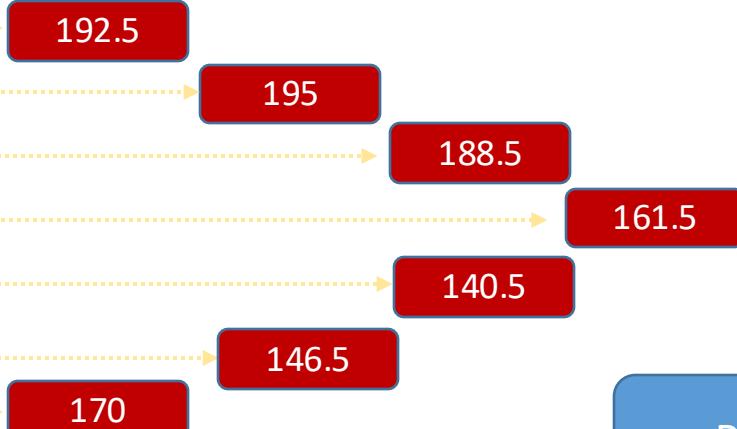
Gini index = $6/8 * (1 - (3/6)^2 - (3/6)^2) + 2/8 * (1 - (1/2)^2 - (1/2)^2) = 0.5$



Gini Index for Heart Disease

Patient Weight	Heart Disease	Sample Weight
205	Yes	1/8
180	Yes	1/8
210	Yes	1/8
167	Yes	1/8
156	No	1/8
125	No	1/8
168	No	1/8
172	No	1/8

$$\text{Gini index} = 4/8 * (1-(1/4)^2 - (3/4)^2) + 4/8 * (1-(1/4)^2 - (3/4)^2) = 0.375$$



PATIENT WEIGHT > 170

HEART DISEASE

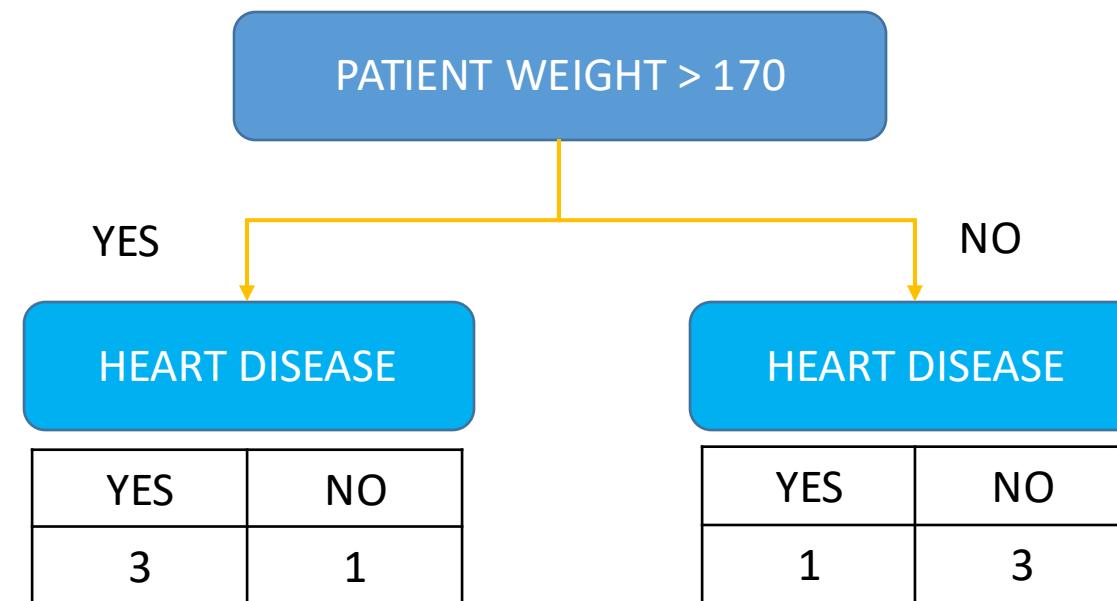
YES	NO
3	1

HEART DISEASE

YES	NO
1	3

Amount of Say

How was this stump contribute to the final decision (classification)?



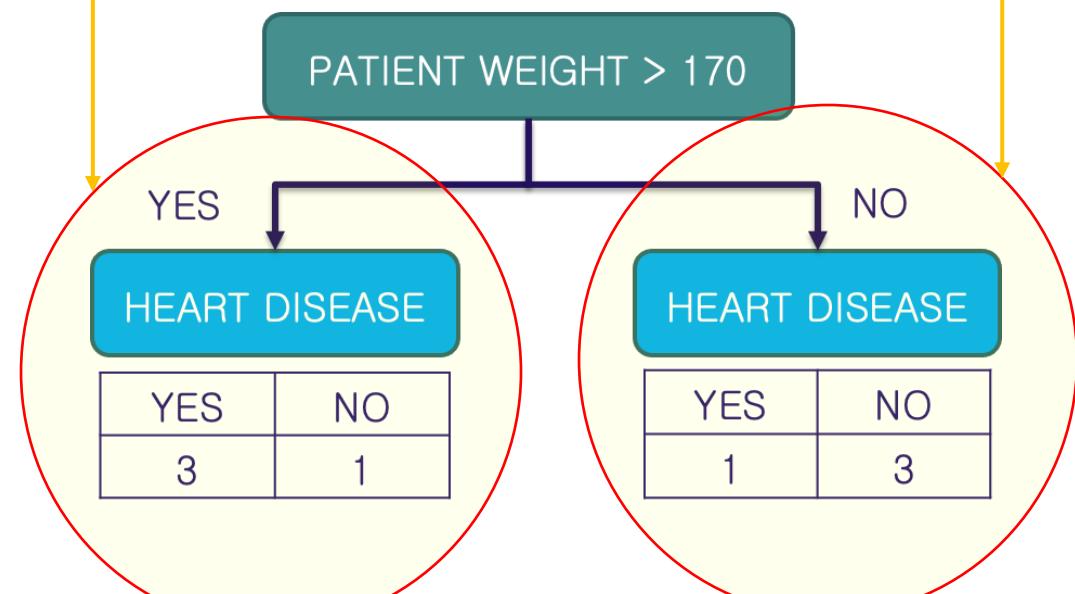
$$\text{Amount of Say} = \frac{1}{2} \log\left(\frac{1 - \text{Total Error}}{\text{Total Error}}\right)$$



Amount of Say: Patient Weight

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8
No	Yes	156	No	1/8
No	Yes	125	No	1/8
Yes	No	168	No	1/8
Yes	Yes	172	No	1/8

- Total error is equal to the sum of the weights of the incorrect classified
- Amount of say = $1/2 * \log((1-2/8) / (2/8)) = 0.55$



Amount of Say: Weight of The Tree

Probability Vs. Odds



Your friend went fishing 10 times a month

- Caught a fish 4 times
- Failed to catch 6 times

What is the *probability* and *odds* of getting a Fish for lunch?

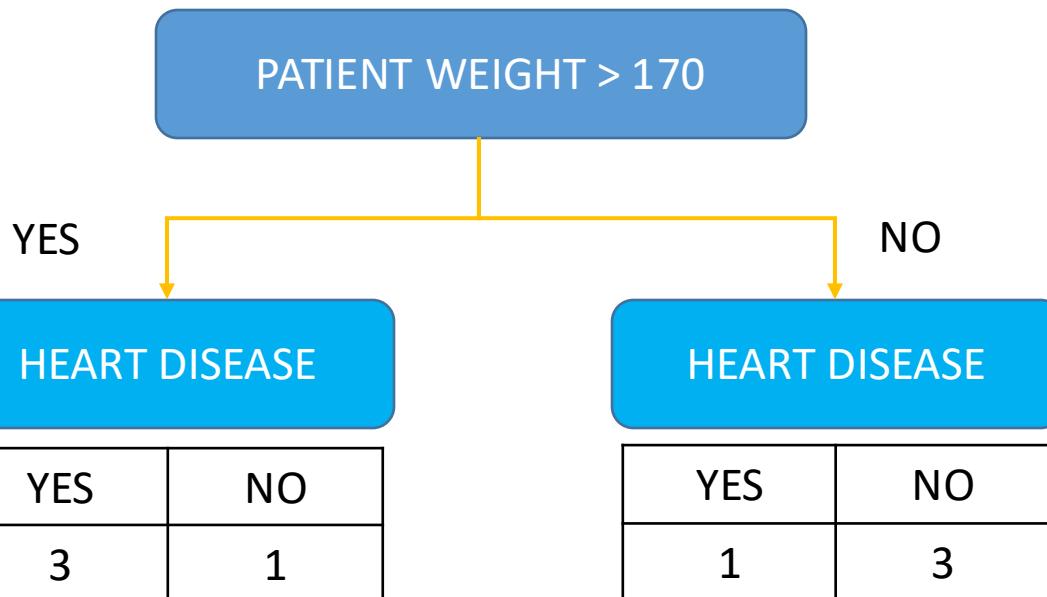
$$\text{Probability} = \frac{\text{Chance for catching fish}}{\text{Total chances}} = \frac{4}{10} = 0.4$$

$$\text{Odds} = \frac{\text{Chance for catching fish}}{\text{Chance for not catching fish}} = \frac{4}{6} = 0.67$$

$$\text{Odds} = \frac{\text{Probability of catching fish}}{\text{Probability of not catching fish}} = \frac{4/10}{6/10} = 0.67$$

$$\text{Odds} = \frac{1 - \text{Probability of not catching fish}}{\text{Probability of not catching fish}} = \frac{4/10}{6/10} = 0.67$$

Amount of Say: Weight of The Tree



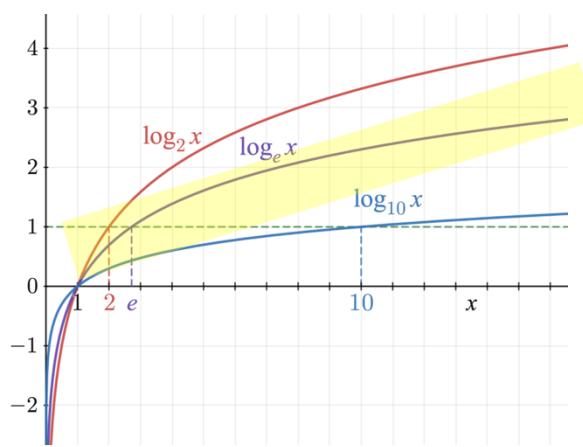
$$\text{Amount of Say} = \frac{1}{2} \log\left(\frac{1 - \text{Total Error}}{\text{Total Error}}\right)$$

$$\text{Odds} = \frac{1 - \text{Probability of not catching fish}}{\text{Probability of not catching fish}} = \frac{4/10}{6/10} = 0.67$$

$$\text{Odds} = \frac{1 - \text{Probability of incorrect prediction}}{\text{Probability of incorrect prediction}}$$

Amount of says = Odds or Amount of says = $\frac{1}{2} \times \log(\text{Odds})$

Why?



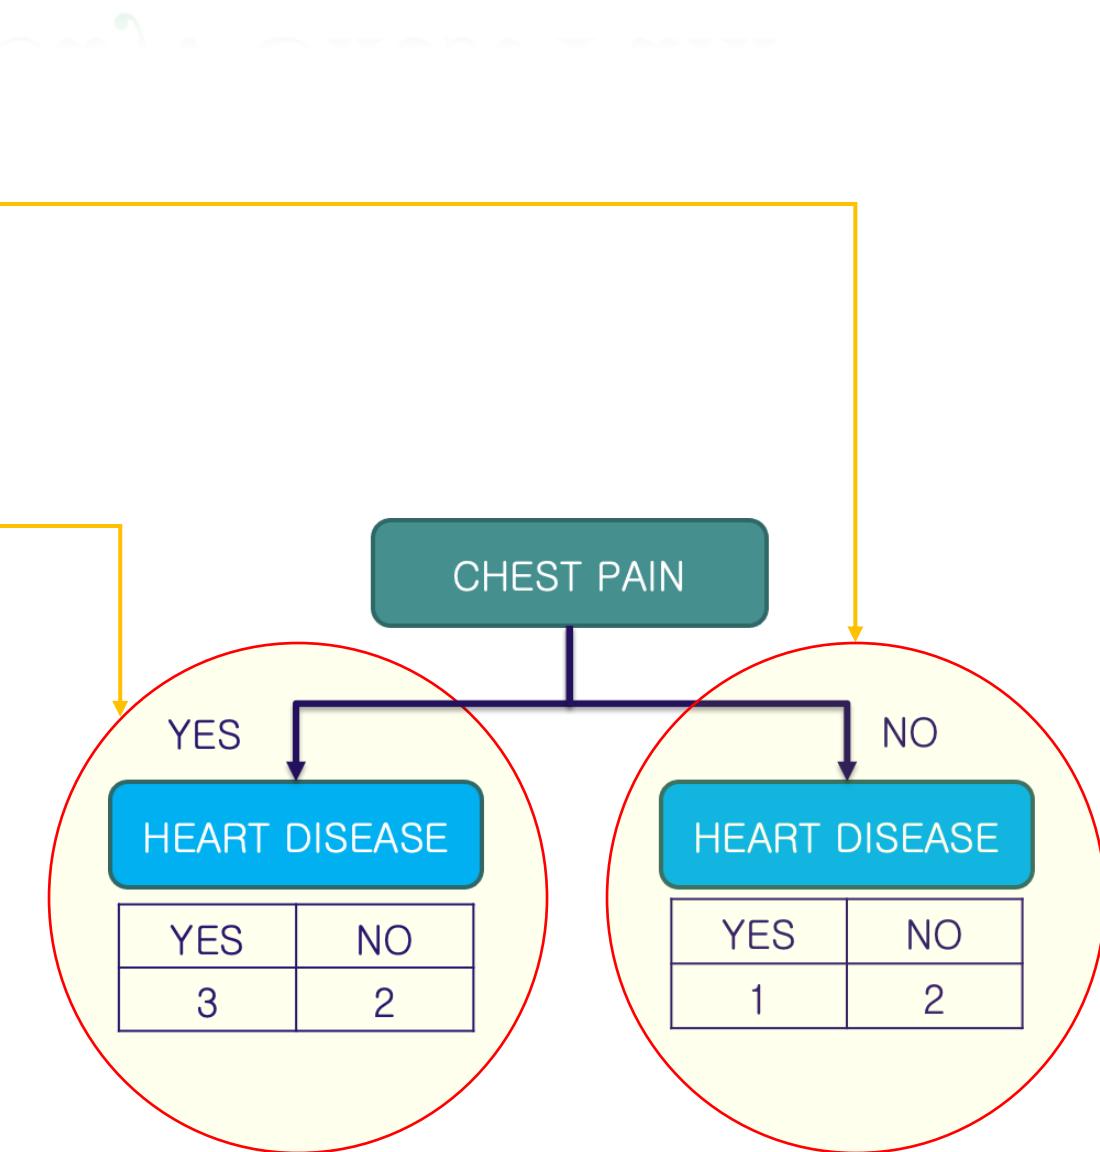
Any Questions

Amount of Say: Chest Pain

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8
No	Yes	156	No	1/8
No	Yes	125	No	1/8
Yes	No	168	No	1/8
Yes	Yes	172	No	1/8

- Total error is equal to the sum of the weights of the incorrect classified
- Amount of say = $1/2 * \log((1 - 3/8) / (3/8)) = 0.25$

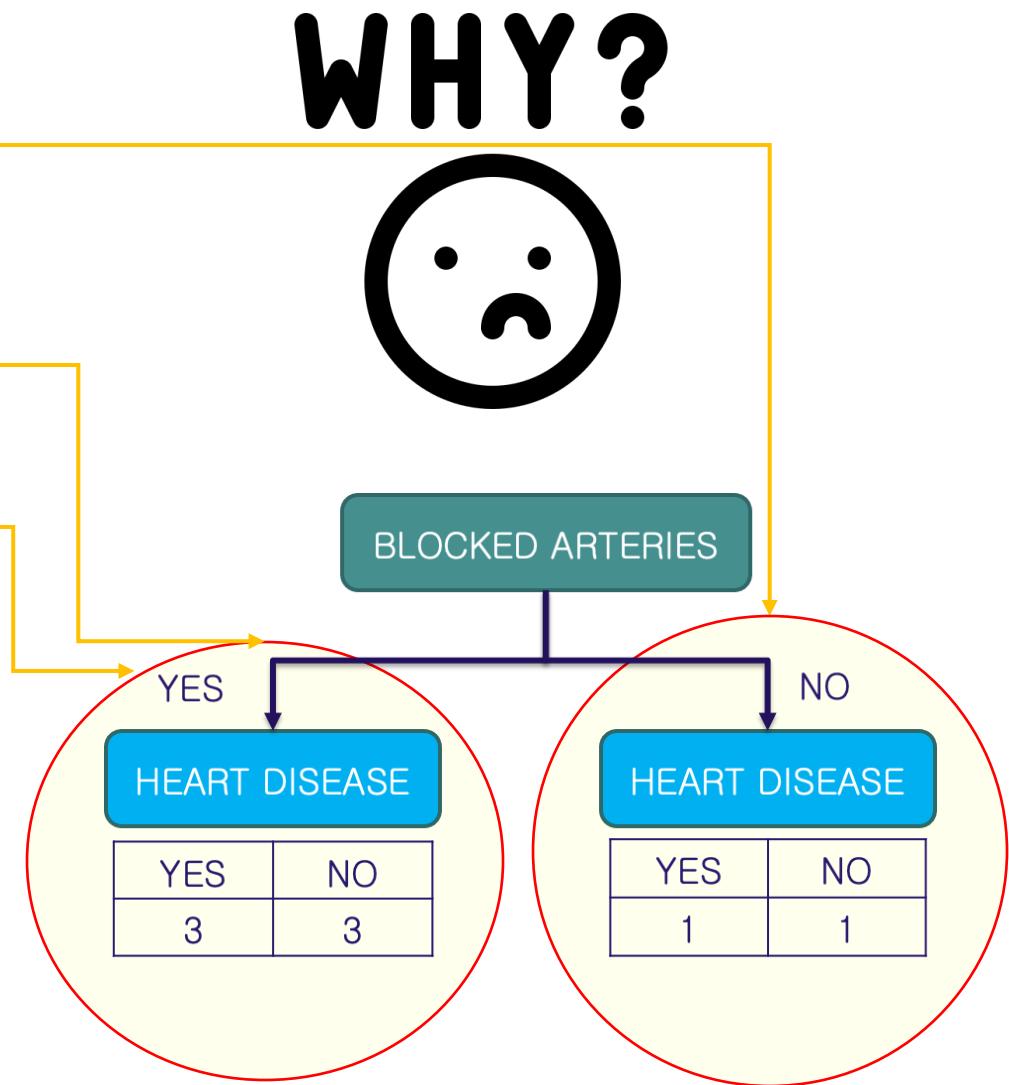
$$\log(\text{Odds}) = \log\left(\frac{1 - \text{Probability of incorrect prediction}}{\text{Probability of incorrect prediction}}\right)$$



Amount of Say: Blocked Arteries

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8
No	Yes	156	No	1/8
No	Yes	125	No	1/8
Yes	No	168	No	1/8
Yes	Yes	172	No	1/8

- Total error is equal to the sum of the weights of the incorrect classified
- Amount of say = $1/2 * \log((1-4/8) / (4/8)) = 0$



Assumptions

Known: weight cho các sample dự đoán sai được sử dụng để tính “Amount of Say” cho từng stump hiện tại.



Unknown: Tiếp theo, chúng ta cần làm thế nào để sử dụng thông tin các weight của sample dự đoán sai này để xây dựng stump tiếp và khắc phục các dự đoán sai này

Idea: Improved Bootstrapped Dataset

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease
Yes	Yes	205	Yes
No	Yes	180	Yes
Yes	No	210	Yes
Yes	Yes	167	Yes
No	Yes	156	No
No	Yes	125	No
Yes	No	168	No
Yes	Yes	172	No

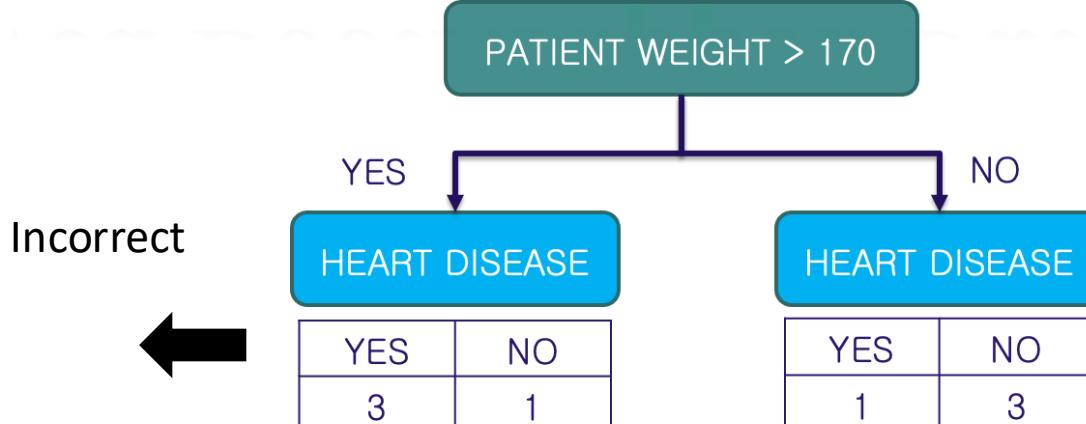
Incorrect

Incorrect



Create new dataset

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease
Yes	Yes	205	Yes
No	Yes	180	Yes
Yes	No	210	Yes
Yes	Yes	167	Yes
Yes	Yes	167	Yes
Yes	Yes	167	Yes
Yes	Yes	172	No
Yes	Yes	172	No



Incorrect

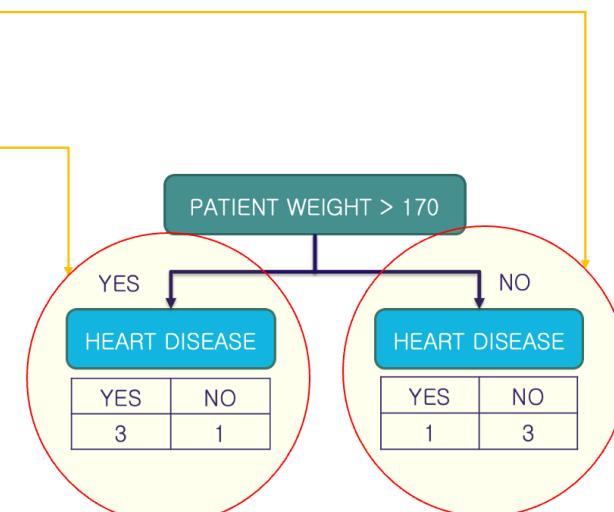
Incorrect



How to Build Next Stump

Increase the sample weights of samples that were incorrectly classified and decrease sample weights of samples that were correctly classified. Label {-1, 1}

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8
No	Yes	156	No	1/8
No	Yes	125	No	1/8
Yes	No	168	No	1/8
Yes	Yes	172	No	1/8



- Total error is equal to the sum of the weights of the incorrect classified
- Amount of say = $1/2 * \log((1-2/8) / (2/8)) = 0.55$

New Sample = sample weight $\times e^{\text{amount of say}}$
Weight

New sample weight = $1/8 * e^{0.55} = 0.22$

```

def update_weights_formular1(w_i, alpha, y, y_pred):
    result = w_i * np.exp(-alpha * y * y_pred)
    w_norm = result / np.sum(result)
    return w_norm
  
```

How to Build Next Stump

Increase the sample weights of samples that were incorrectly classified and decrease the sample weights of samples that were correctly classified. Label {-1, 1}

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease
Yes	Yes	205	Yes
No	Yes	180	Yes
Yes	No	210	Yes
Yes	Yes	167	Yes
No	Yes	156	No
No	Yes	125	No
Yes	No	168	No
Yes	Yes	172	No

New Sample Weight = sample weight $\times e^{-\text{amount of say}}$

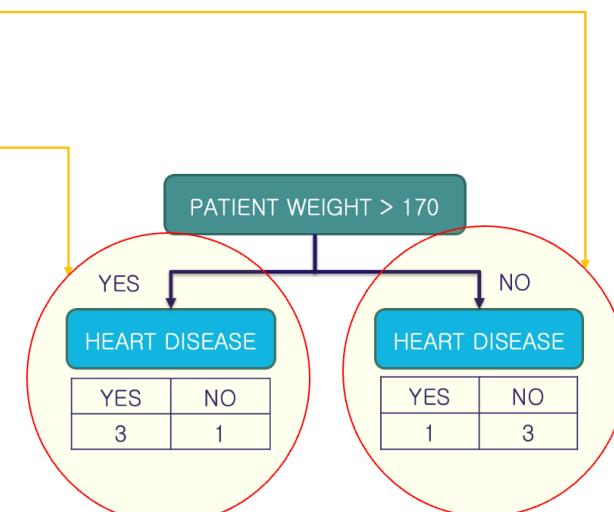
New sample weight = $1/8 * e^{-0.55} = 0.07$

```
def update_weights_formular1(w_i, alpha, y, y_pred):
    result = w_i * np.exp(-alpha * y * y_pred)
    w_norm = result / np.sum(result)
    return w_norm
```

How to Build Next Stump

Increase the sample weights of samples that were incorrectly classified and keep the sample weights of samples that were correctly classified. Label {0, 1}

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8
No	Yes	156	No	1/8
No	Yes	125	No	1/8
Yes	No	168	No	1/8
Yes	Yes	172	No	1/8



- Total error is equal to the sum of the weights of the incorrect classified
- Amount of say = $1/2 * \log((1-2/8) / (2/8)) = 0.55$

$$w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))],$$

New sample weight = $1/8 * e^{0.55*1} = 0.22$

```

def update_weights_formular2(w_i, alpha, y, y_pred):
    result = w_i * np.exp(alpha * (
        np.not_equal(y, y_pred)).astype(int))
    w_norm = result / np.sum(result)
    return w_norm
  
```

How to Build Next Stump

Increase the sample weights of samples that were incorrectly classified and **keep the sample weights of samples that were correctly classified**. Label {0, 1}

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease
Yes	Yes	205	Yes
No	Yes	180	Yes
Yes	No	210	Yes
Yes	Yes	167	Yes
No	Yes	156	No
No	Yes	125	No
Yes	No	168	No
Yes	Yes	172	No

$$w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$$

New sample weight = $1/8 * e^{0.55*0} = 0.125$

```
def update_weights_formular2(w_i, alpha, y, y_pred):
    result = w_i * np.exp(alpha * (
        np.not_equal(y, y_pred).astype(int)))
    w_norm = result / np.sum(result)
    return w_norm
```

New Sample Weight

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight	New Weight	Normal Weight
Yes	Yes	205	Yes	1/8	0.07	0.08
No	Yes	180	Yes	1/8	0.07	0.08
Yes	No	210	Yes	1/8	0.07	0.08
Yes	Yes	167	Yes	1/8	0.22	0.25
No	Yes	156	No	1/8	0.07	0.08
No	Yes	125	No	1/8	0.07	0.08
Yes	No	168	No	1/8	0.07	0.08
Yes	Yes	172	No	1/8	0.22	0.25
Sum				~1.0	0.86	~1.0

New Sample Weight = sample weight $\times e^{\text{amount of say}}$

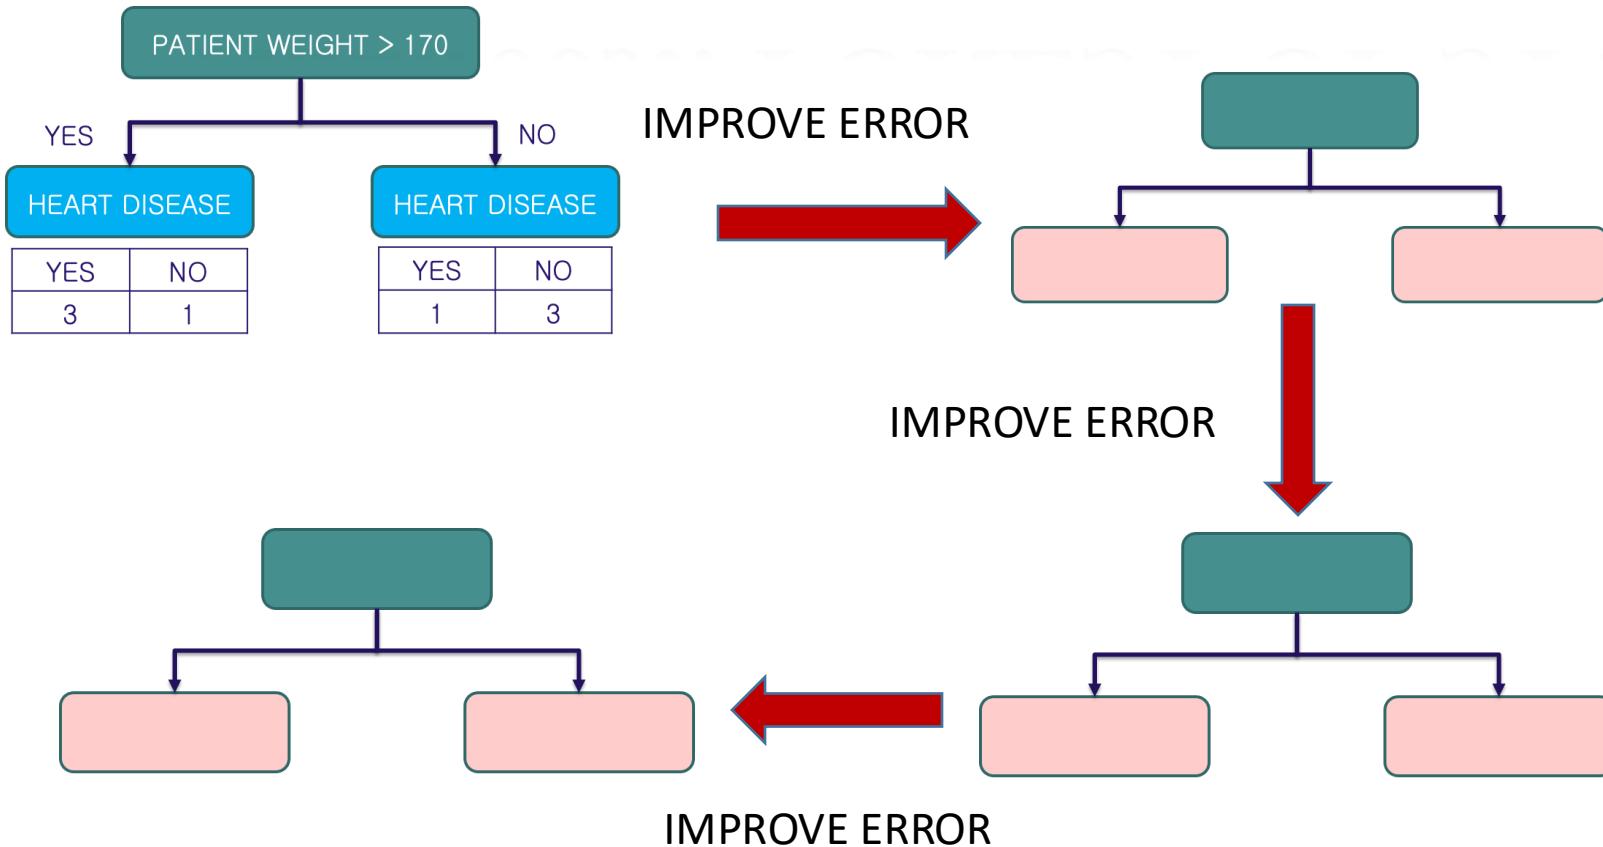


Update

New Sample Weight

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	New Weight
Yes	Yes	205	Yes	0.08
No	Yes	180	Yes	0.08
Yes	No	210	Yes	0.08
Yes	Yes	167	Yes	0.25
No	Yes	156	No	0.08
No	Yes	125	No	0.08
Yes	No	168	No	0.08
Yes	Yes	172	No	0.25
Sum				~1.0

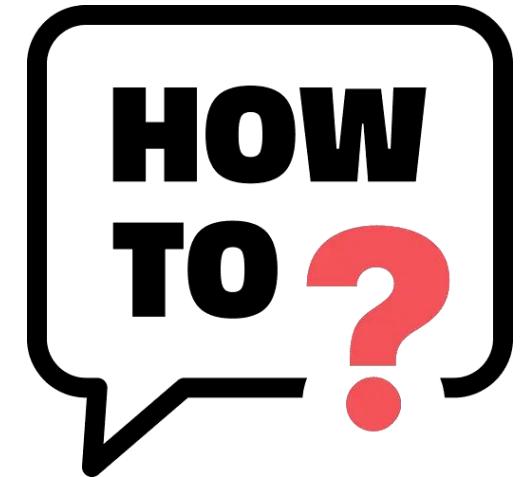
AdaBoost: FOREST OF STUMPS



?
HOW

New Dataset

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Normal Weight
Yes	Yes	205	Yes	0.08
No	Yes	180	Yes	0.08
Yes	No	210	Yes	0.08
Yes	Yes	167	Yes	0.25
No	Yes	156	No	0.08
No	Yes	125	No	0.08
Yes	No	168	No	0.08
Yes	Yes	172	No	0.25
Sum			~ 1.0	

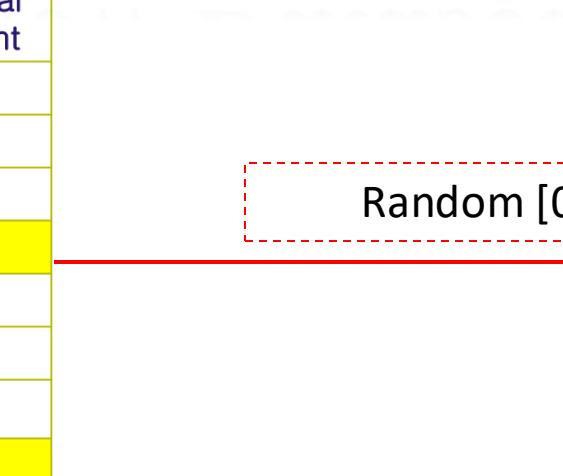


Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Normal Weight
Yes	Yes	205	Yes	0.08
No	Yes	180	Yes	0.08
Yes	No	210	Yes	0.08
Yes	Yes	167	Yes	0.25
No	Yes	156	No	0.08
No	Yes	125	No	0.08
Yes	No	168	No	0.08
Yes	Yes	172	No	0.25
Sum			~ 1.0	

New Dataset

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Normal Weight
Yes	Yes	205	Yes	0.08
No	Yes	180	Yes	0.08
Yes	No	210	Yes	0.08
Yes	Yes	167	Yes	0.25
No	Yes	156	No	0.08
No	Yes	125	No	0.08
Yes	No	168	No	0.08
Yes	Yes	172	No	0.25
Sum			~ 1.0	

Random [0,1] to select samples



Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Normal Weight
Yes	Yes	205	Yes	0.08
No	Yes	180	Yes	0.08
Yes	No	210	Yes	0.08
Yes	Yes	167	Yes	0.25
No	Yes	156	No	0.08
No	Yes	125	No	0.08
Yes	No	168	No	0.08
Yes	Yes	172	No	0.25
Sum			~ 1.0	

New Dataset

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Normal Weight	Range
Yes	Yes	205	Yes	0.08	[0-0.08]
No	Yes	180	Yes	0.08	(0.08-0.16]
Yes	No	210	Yes	0.08	(0.16-0.24]
Yes	Yes	167	Yes	0.25	(0.24-0.495]
No	Yes	156	No	0.08	(0.495-0.575]
No	Yes	125	No	0.08	(0.575-0.655]
Yes	No	168	No	0.08	(0.655-0.735]
Yes	Yes	172	No	0.25	(0.735-1.0]
Sum				~1.0	

New Dataset

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Normal Weight
Yes	Yes	205	Yes	0.08
No	Yes	180	Yes	0.08
Yes	No	210	Yes	0.08
Yes	Yes	167	Yes	0.25
No	Yes	156	No	0.08
No	Yes	125	No	0.08
Yes	No	168	No	0.08
Yes	Yes	172	No	0.25
Sum				~1.0

Old dataset

Ý tưởng: các sample bị phân loại sai, sẽ được thêm vào dataset mới

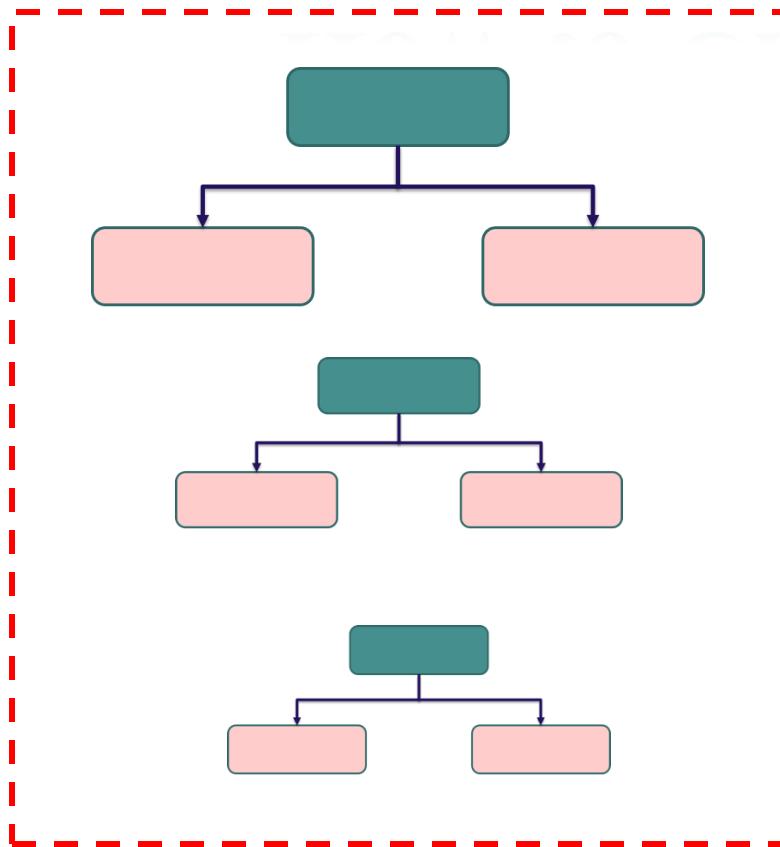
Random [0,1] to select samples

CONTINUE TO BUILD THE NEXT STUMP

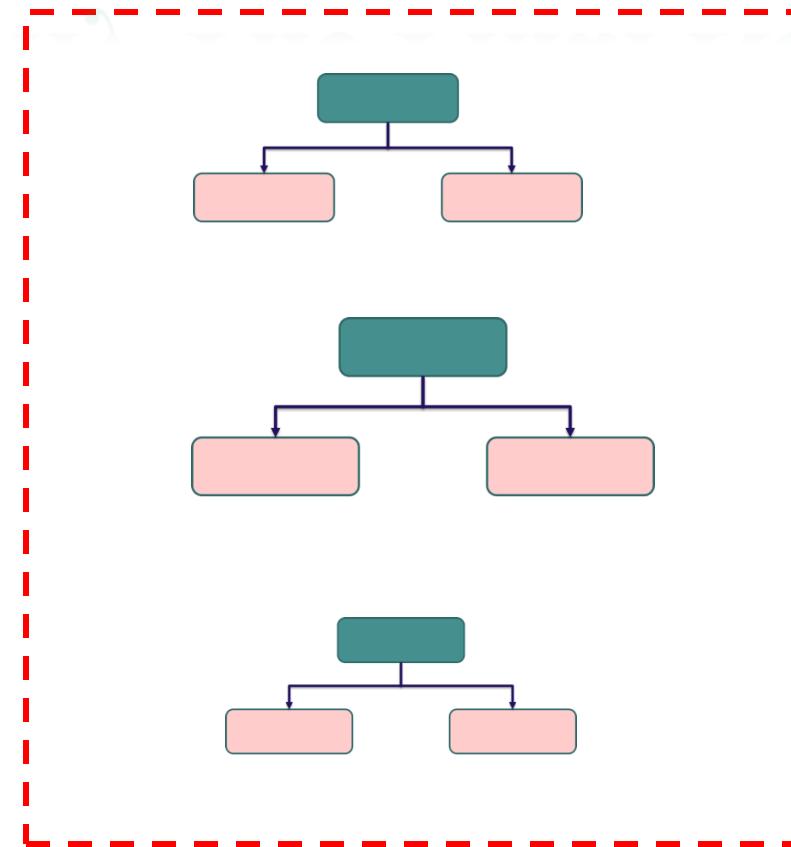
Chest Pain	Blocked Arteries	Patient Weight	Heart Disease
Yes	Yes	167	Yes
No	Yes	180	Yes
Yes	No	210	Yes
Yes	Yes	167	Yes
No	Yes	156	No
No	Yes	125	No
Yes	Yes	172	No
Yes	Yes	172	No

New dataset

How to Classify The Final Result



These stumps for predicting
heart disease



These stumps for predicting
no heart disease

How to Classify The Final Result

Algorithm 1 AdaBoost (*Freund & Schapire 1997*)

1. Initialize the observation weights $w_i = 1/n$, $i = 1, 2, \dots, n$.

2. For $m = 1$ to M :

(a) Fit a classifier $T^{(m)}(\mathbf{x})$ to the training data using weights w_i .

(b) Compute

$$err^{(m)} = \sum_{i=1}^n w_i \mathbb{I}(c_i \neq T^{(m)}(\mathbf{x}_i)) / \sum_{i=1}^n w_i.$$

(c) Compute

$$\alpha^{(m)} = \log \frac{1 - err^{(m)}}{err^{(m)}}.$$

(d) Set

$$w_i \leftarrow w_i \cdot \exp \left(\alpha^{(m)} \cdot \mathbb{I}(c_i \neq T^{(m)}(\mathbf{x}_i)) \right), \quad i = 1, 2, \dots, n.$$

(e) Re-normalize w_i .

3. Output

$$C(\mathbf{x}) = \arg \max_k \sum_{m=1}^M \alpha^{(m)} \cdot \mathbb{I}(T^{(m)}(\mathbf{x}) = k).$$

How to Classify The Final Result

SAMME — Stagewise
Additive Modeling using a
Multi-class Exponential
loss function

Algorithm 2 SAMME

1. Initialize the observation weights $w_i = 1/n, i = 1, 2, \dots, n.$
2. For $m = 1$ to M :
 - (a) Fit a classifier $T^{(m)}(\mathbf{x})$ to the training data using weights w_i .
 - (b) Compute

$$\text{err}^{(m)} = \sum_{i=1}^n w_i \mathbb{I}(c_i \neq T^{(m)}(\mathbf{x}_i)) / \sum_{i=1}^n w_i.$$

- (c) Compute

$$\alpha^{(m)} = \log \frac{1 - \text{err}^{(m)}}{\text{err}^{(m)}} + \log(K - 1). \quad (1)$$

- (d) Set

$$w_i \leftarrow w_i \cdot \exp \left(\alpha^{(m)} \cdot \mathbb{I}(c_i \neq T^{(m)}(\mathbf{x}_i)) \right), \quad i = 1, \dots, n.$$

- (e) Re-normalize w_i .

3. Output

$$C(\mathbf{x}) = \arg \max_k \sum_{m=1}^M \alpha^{(m)} \cdot \mathbb{I}(T^{(m)}(\mathbf{x}) = k).$$

Behind The Scene

The hypothesis function $f(x)$ is defined as:

$$f(x) = \sum_{m=1}^M \alpha_m * G_m(x)$$

M trees Label: {-1,1} N Samples

$$Err(f) = \sum_{i=1}^N e^{-y_i * f(x_i)}$$

Exponential loss function:

$$\boxed{f_m(x) = f_{m-1}(x) + \alpha_m * G_m(x)}$$

Supposing that:

$$\begin{aligned} Err(f) &= \sum_{i=1}^N e^{-y_i f_{m-1}(x_i)} * e^{-y_i \alpha_m G_m(x_i)} \\ \Rightarrow Err(\alpha_m, G_m) &= \sum_{i=1}^N e^{-y_i f_{m-1}(x_i)} * e^{-y_i \alpha_m G_m(x_i)} \end{aligned}$$

$w_i = e^{-y_i f_{m-1}(x_i)}$
not dependent
on
 α_m and G_m

Behind The Scene

Expand the error function:

$$Err(f) = \sum_{i=1}^N e^{-y_i f_{m-1}(x_i)} * e^{-y_i \alpha_m G_m(x_i)}$$

$$\Rightarrow Err(\alpha_m, G_m) = \sum_{i=1}^N e^{-y_i f_{m-1}(x_i)} * e^{-y_i \alpha_m G_m(x_i)}$$

$$w_i = e^{-y_i f_{m-1}(x_i)}$$

$$Err = e^{-\alpha_m} \sum_{y_i=G_m(x_i)} w_i + e^{\alpha_m} \sum_{y_i \neq G_m(x_i)} w_i$$

It's easy to show that the expression $-y_i \alpha_m G_m(x_i)$ is $-\alpha_m$ if $y_i = G_m(x_i)$
and is α_m if $y_i \neq G_m(x_i)$.

Total

weight $T_w = \sum w_i$

Error weight E_w :

$$E_w = \sum_{y_i \neq G_m(x_i)} w_i$$

$$Err = e^{-\alpha_m} \sum_{y_i=G_m(x_i)} w_i + e^{\alpha_m} \sum_{y_i \neq G_m(x_i)} w_i$$

$$Err = e^{-\alpha_m} (T_w - E_w) + e^{\alpha_m} E_w$$

$$\frac{dErr}{d\alpha_m} = -\alpha_m e^{-\alpha_m} (T_w - E_w) + \alpha_m e^{\alpha_m} E_w$$

M trees

N Samples

Label: {-1,1}

Behind The Scene

$$Err = e^{-\alpha_m} (T_w - E_w) + e^{\alpha_m} E_w$$

$$\frac{dErr}{d\alpha_m} = -\alpha_m e^{-\alpha_m} (T_w - E_w) + \alpha_m e^{\alpha_m} E_w$$

$$\begin{aligned} -\alpha_m e^{-\alpha_m} (T_w - E_w) + \alpha_m e^{\alpha_m} E_w &= 0 \\ \Rightarrow e^{\alpha_m} E_w - e^{-\alpha_m} (T_w - E_w) &= 0 \\ \Rightarrow e^{\alpha_m} E_w &= e^{-\alpha_m} (T_w - E_w) \\ \Rightarrow e^{2\alpha_m} &= (T_w - E_w)/E_w \\ \Rightarrow e^{2\alpha_m} &= \frac{1 - (E_w/T_w)}{(E_w/T_w)} \end{aligned}$$

Taking log on both sides we have

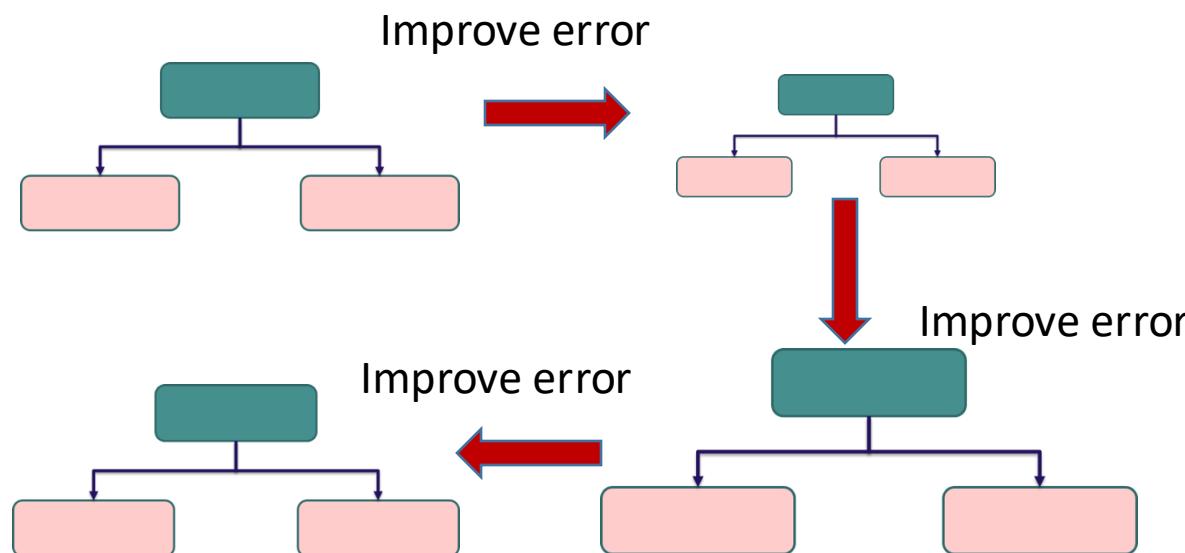


$$\alpha_m = \frac{1}{2} \log \frac{1 - err_m}{err_m}$$

$$err_m = \frac{E_w}{T_w} = \frac{\sum_{y_i \neq G_m(x_i)} w_i}{\sum w_i}$$

Summary

Height	Favorite Color	Gender	Weight
1.6	Blue	Male	88
1.6	Green	Female	76
1.5	Blue	Female	56
1.8	Red	Male	73
1.5	Green	Male	77
1.4	Blue	Female	57



AdaBoost builds a stump based on the error made by previous stumps

Example: Spam Classification

- Classifying Email as Spam or Non-Spam

 Spambase Donated on 6/30/1999		
Classifying Email as Spam or Non-Spam		
Dataset Characteristics	Subject Area	Associated Tasks
Multivariate	Computer Science	Classification
Attribute Type	# Instances	# Attributes
Integer, Real	4601	57

<http://archive.ics.uci.edu/dataset/94/spambase>

Example: Spam Classification

1. Initialize the observation weights $w_i = 1/N, i = 1, 2, \dots, N$.
2. For $m = 1$ to M :
 - (a) Fit a classifier $G_m(x)$ to the training data using weights w_i .
 - (b) Compute
$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$
 - (c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.
 - (d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))], i = 1, 2, \dots, N$.
3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.

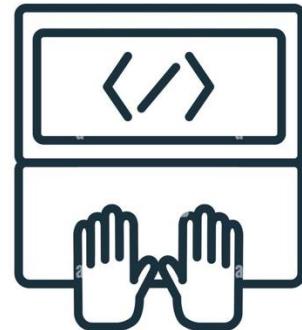
Example: Spam Classification

Our implementation

```
# Fit model
ab = AIVNAdaBoost()
ab.fit(X_train, y_train, M = 50)

# Predict on test set
y_pred = ab.predict(X_test)
print('The accuracy_score of the model is:', round(accuracy_score(y_test, y_pred), 4))
```

The accuracy_score of the model is: 0.9392



CUSTOM CODE

```
▶ from sklearn.ensemble import AdaBoostClassifier
|
ab_sk = AdaBoostClassifier(n_estimators = 50)
ab_sk.fit(X_train, y_train)
y_pred_sk = ab_sk.predict(X_test)
print('The accuracy_score of the model is:', round(accuracy_score(y_test, y_pred_sk), 4))
```

The accuracy_score of the model is: 0.9435

Using Sklearn library

Outline



- **Boosting Techniques**
- **AdaBoost Clearly Explain**
- **Gradient Boost Clearly Explain**
- **Time Series Data: Predicting Energy Consumption**
- **Summary**

Gradient Boost For Regression

Height	Favorite Color	Gender	Weight
1.6	Blue	Male	88
1.6	Green	Female	76
1.5	Blue	Female	56
1.8	Red	Male	73
1.5	Green	Male	77
1.4	Blue	Female	57



Input



Output

Tree-based Gradient Boost

- Step 1: Build 1st tree
 - Calculate the average of weights

Height	Favorite Color	Gender	Weight
1.6	Blue	Male	88
1.6	Green	Female	76
1.5	Blue	Female	56
1.8	Red	Male	73
1.5	Green	Male	77
1.4	Blue	Female	57

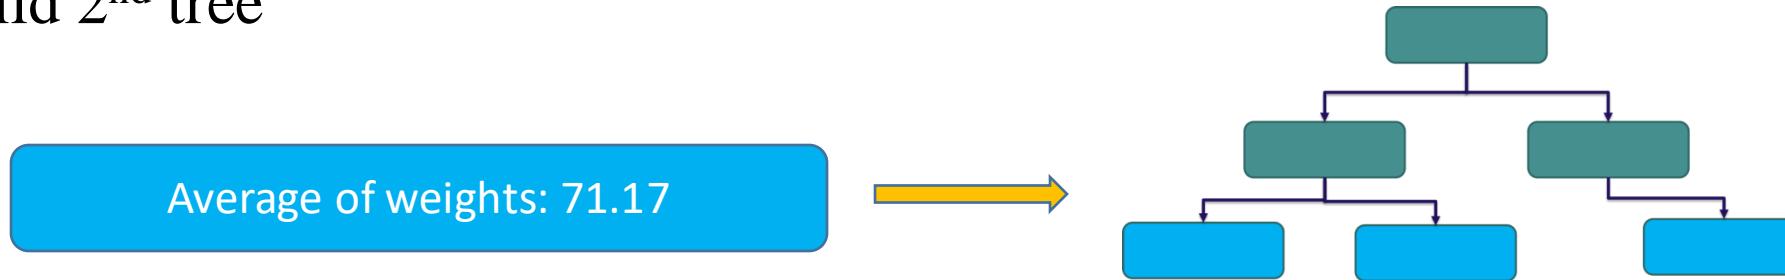
Node of 1st Tree

Average of Weights: 71.17



Tree-based Gradient Boost

- Step 2:
 - Build 2nd tree



Height	Favorite Color	Gender	Weight
1.6	Blue	Male	88
1.6	Green	Female	76
1.5	Blue	Female	56
1.8	Red	Male	73
1.5	Green	Male	77
1.4	Blue	Female	57

Tree-based Gradient Boost

- Step 2:
➤ Build 2nd tree



Height	Favorite Color	Gender	Weight	Residual Error
1.6	Blue	Male	88	16.8
1.6	Green	Female	76	1.8
1.5	Blue	Female	56	-15.2
1.8	Red	Male	73	1.8
1.5	Green	Male	77	5.8
1.4	Blue	Female	57	-14.2

Tree-based Gradient Boost

Height	Favorite Color	Gender	Residual Error
1.6	Blue	Male	16.8
1.6	Green	Female	1.8
1.5	Blue	Female	-15.2
1.8	Red	Male	1.8
1.5	Green	Male	5.8
1.4	Blue	Female	-14.2

Tại sao lại xây dựng cây dự đoán Residual Error

Gender is Female

Height < 1.6

Color is not Blue

-14.2, -15.2

4.8

1.5, 5.8

16.8

Tree-based Gradient Boost

Height	Favorite Color	Gender	Residual Error
1.6	Blue	Male	16.8
1.6	Green	Female	1.8
1.5	Blue	Female	-15.2
1.8	Red	Male	1.8
1.5	Green	Male	5.8
1.4	Blue	Female	-14.2

Tại sao lại xây dựng cây dự đoán Residual Error

Gender is Female

Height < 1.6

Color is not Blue

Trung bình residual

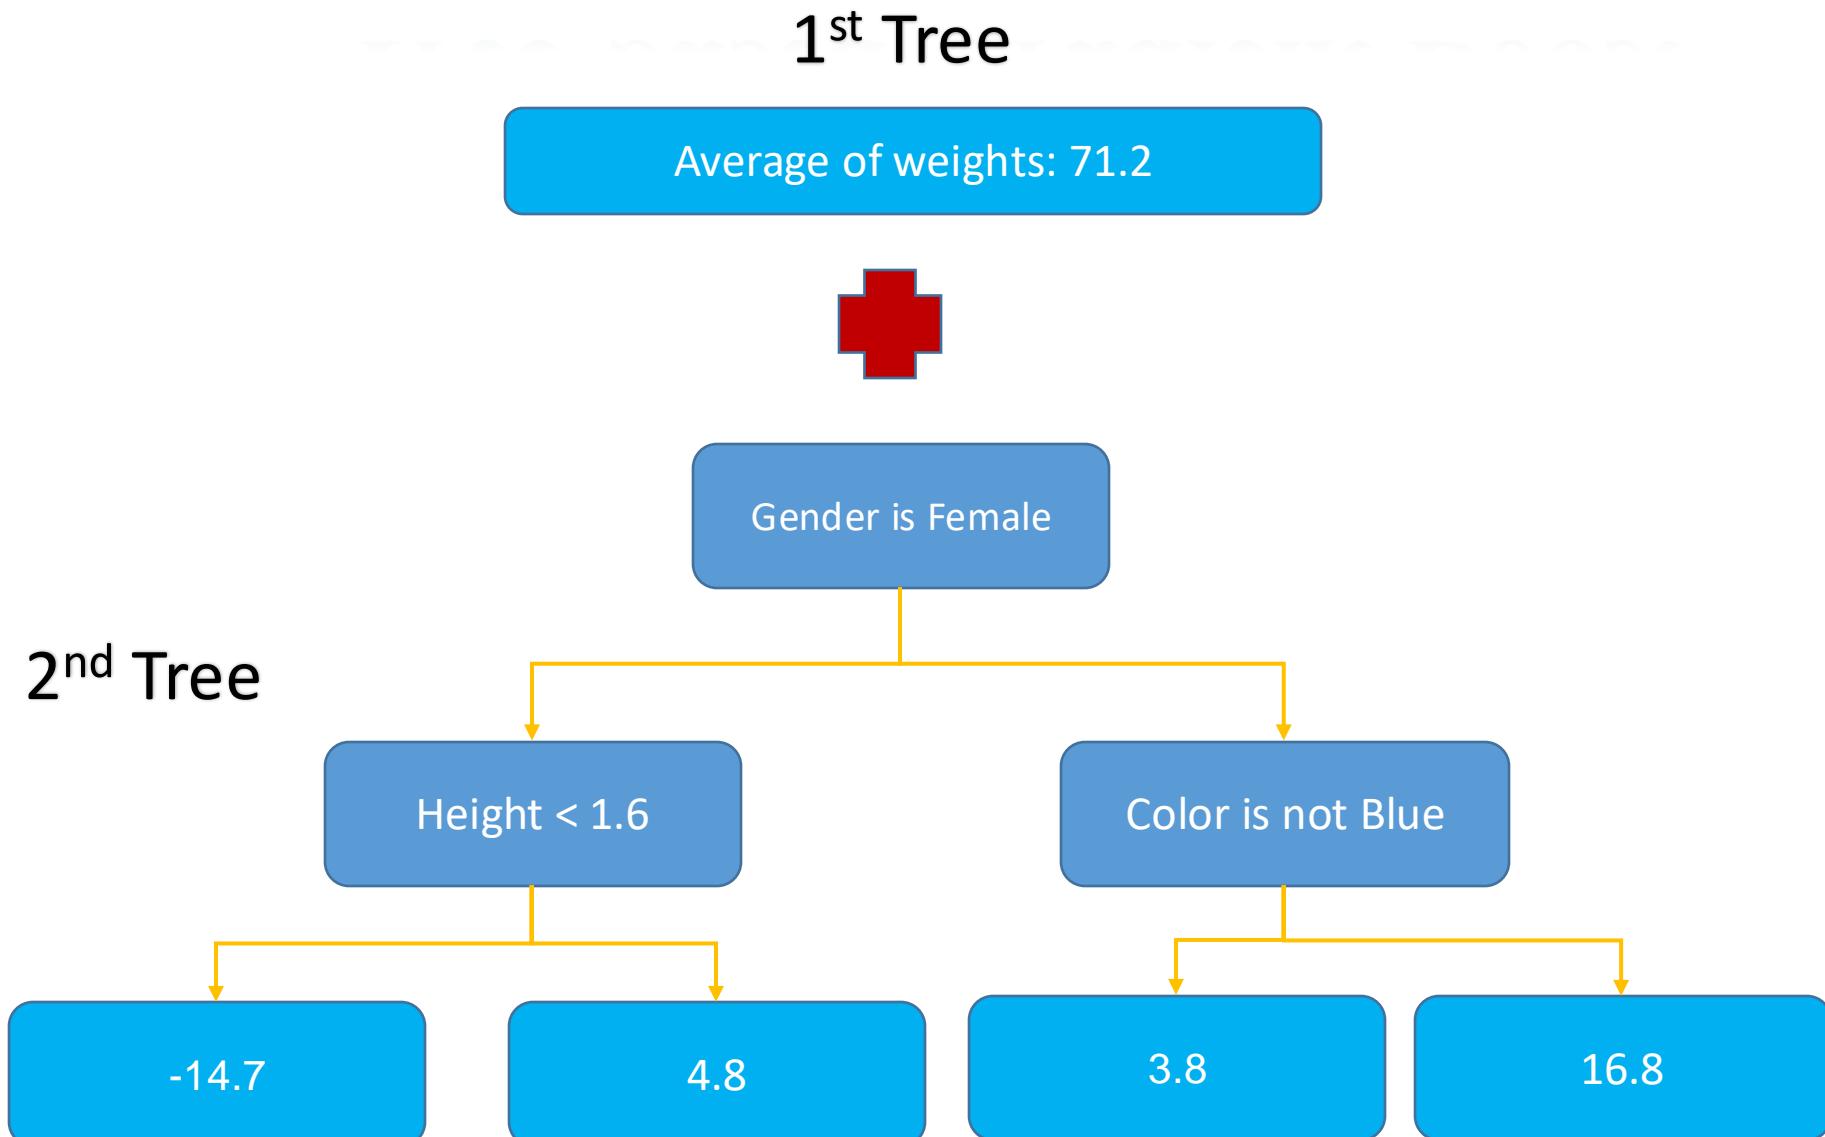
-14.7

4.8

3.8

16.8

Tree-based Gradient Boost



Prediction

Height	Favorite Color	Gender	Weight	Prediction
1.6	Blue	Male	88	88
1.6	Green	Female	76	76
1.5	Blue	Female	56	56
1.8	Red	Male	73	73
1.5	Green	Male	77	77
1.4	Blue	Female	57	57

OVER FITTING

1st Tree

Average of weights: 71.2



Gender is Female

2nd Tree

Height < 1.6

-14.7

Color is not Blue

4.8

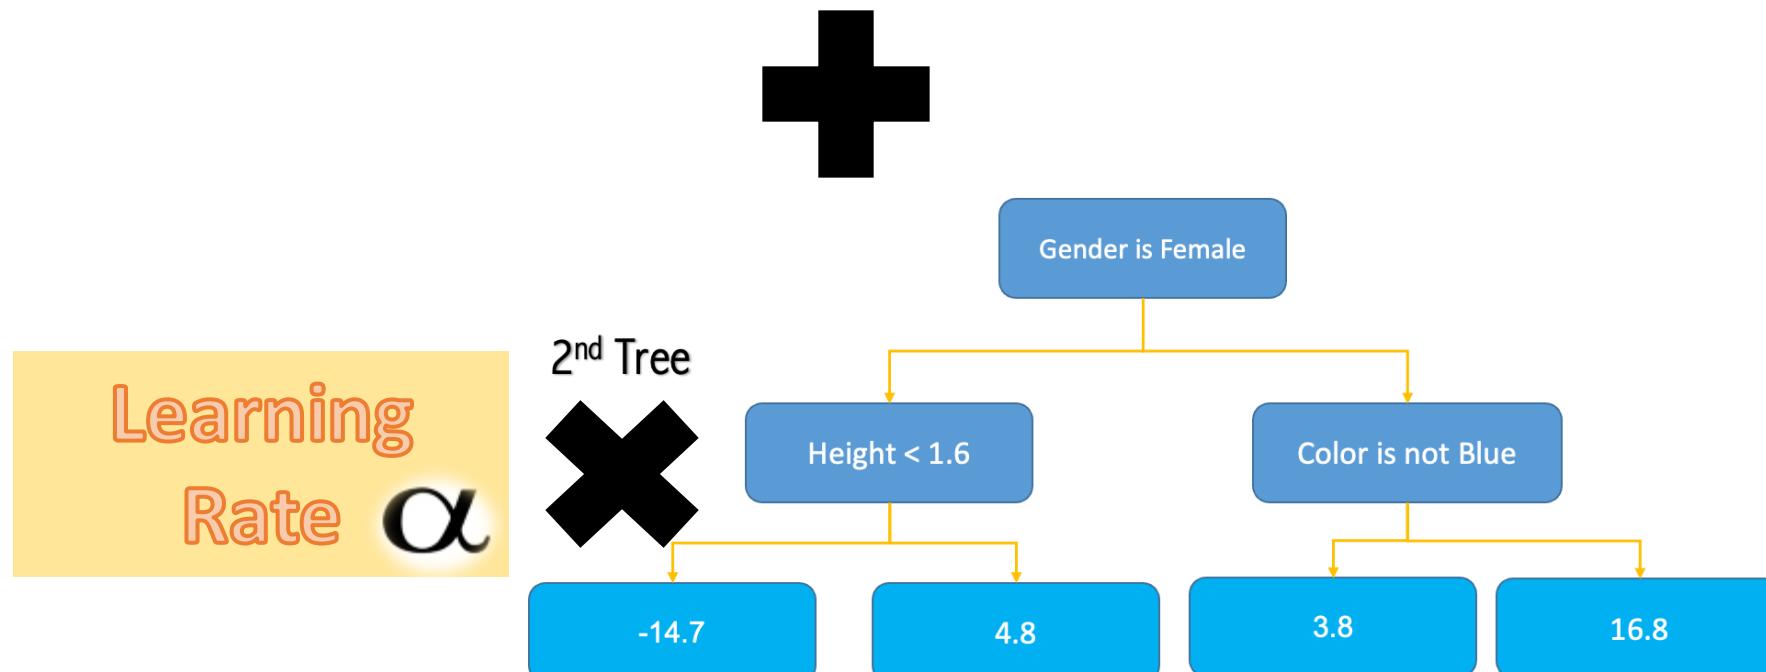
3.8

16.8

Prediction

1st Tree

AVG of weights: 71.2



Prediction

Height	Favorite Color	Gender	Weight	Prediction
1.6	Blue	Male	88	74.56

$$\text{Prediction} = 71.2 + 0.2 * 16.8 = 74.56$$

1st Tree

AVG of weights: 71.2



Gender is Female

2nd Tree



Height < 1.6

-14.7

4.8

Color is not Blue

3.8

16.8

$\alpha = 0.2$

Learning Rate α

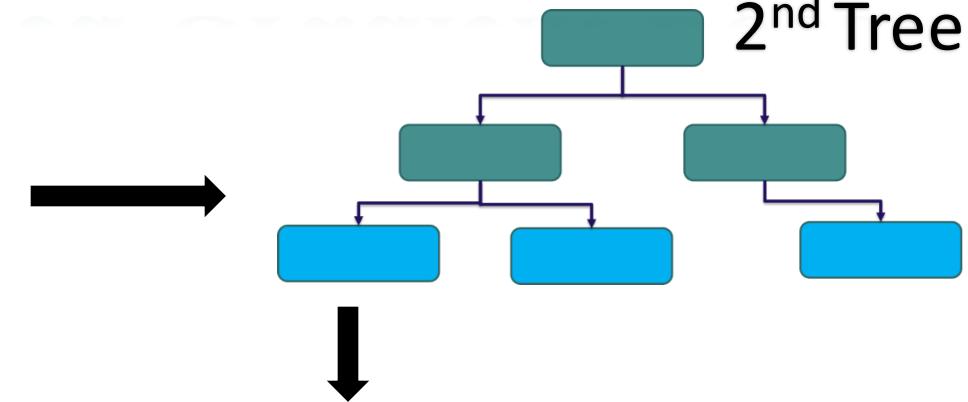
Tree-based Gradient Boost

- Step 3:

➤ Build 3rd tree

1st Tree

Average of weights: 71.2



Height	Favorite Color	Gender	Weight	Predicted Weight	Residual Error
1.6	Blue	Male	88	74.56	12.44
1.6	Green	Female	76
1.5	Blue	Female	56
1.8	Red	Male	73
1.5	Green	Male	77
1.4	Blue	Female	57

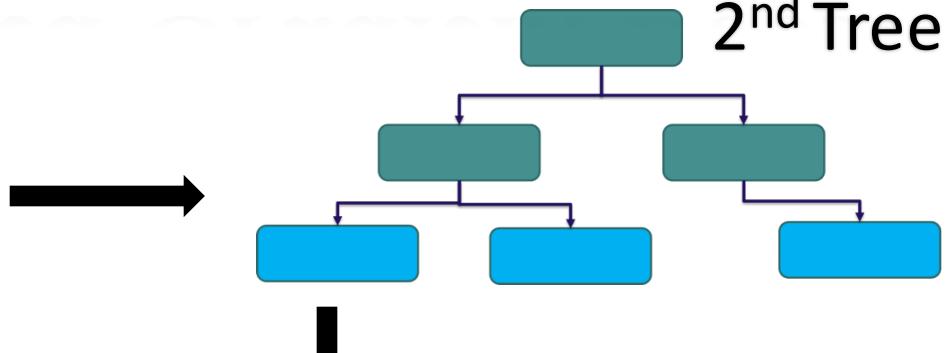
Tree-based Gradient Boost

- Step 3:

➤ Build 3rd tree

1st Tree

Average of weights: 71.2

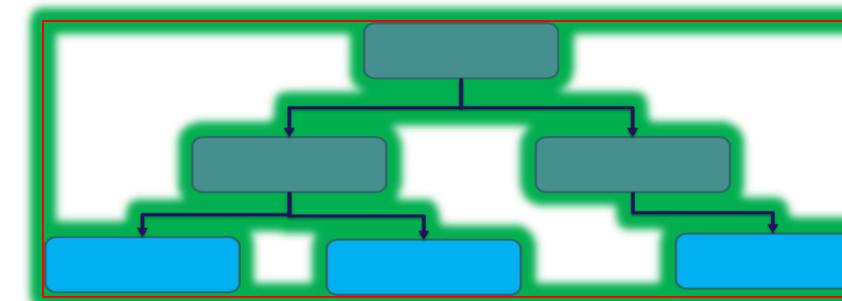
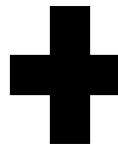
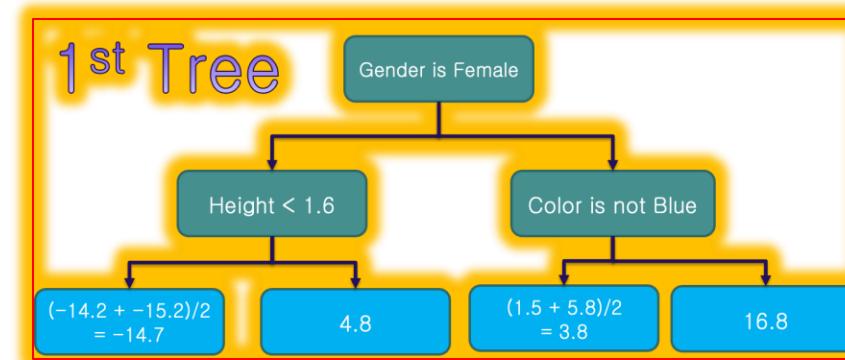
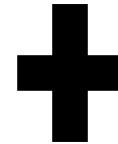


Height	Favorite Color	Gender	First Tree Residual	Second Tree Residual	Third Tree Residual
1.6	Blue	Male	16.8	12.44	
1.6	Green	Female	1.8	...	
1.5	Blue	Female	-15.2	...	
1.8	Red	Male	1.8	...	
1.5	Green	Male	5.8	...	
1.4	Blue	Female	-14.2	...	

Prediction

$\alpha \times \times$

AVG of weights: 71.2



Gradient Boost: Behind The Scenes

**Loss
Function:**



	Height	Favorite Color	Gender	Weight
1.6	Blue	Male	88	
1.6	Green	Female	76	
1.5	Blue	Female	56	
1.8	Red	Male	73	
1.5	Green	Male	77	
1.4	Blue	Female	57	

- $\{(x_i, y_i)\}_{i=1}^n$

- Loss function = $L(y_i, F(x)) = 1/2 * (\text{Output} - \text{Predicted})^2$

$$J = \frac{1}{2m} \sum_{i=1}^m (\hat{y} - y)^2$$

$$\frac{dL}{d\text{Predicted}} = 2/2 (\text{Output} - \text{Predicted}) * -1 = -(\text{Output} - \text{Predicted})$$

Tricky implementation here



CUSTOM CODE

Gradient Boost: Behind The Scenes

- Step 1:
 - Initialize a model with a constant value:
 - $F_0(x) = \underset{\delta}{\operatorname{argmin}} \sum_{i=1}^n L(y_i, \delta)$

Height	Favorite Color	Gender	Weight	y
1.6	Blue	Male	88	
1.6	Green	Female	76	
1.5	Blue	Female	56	

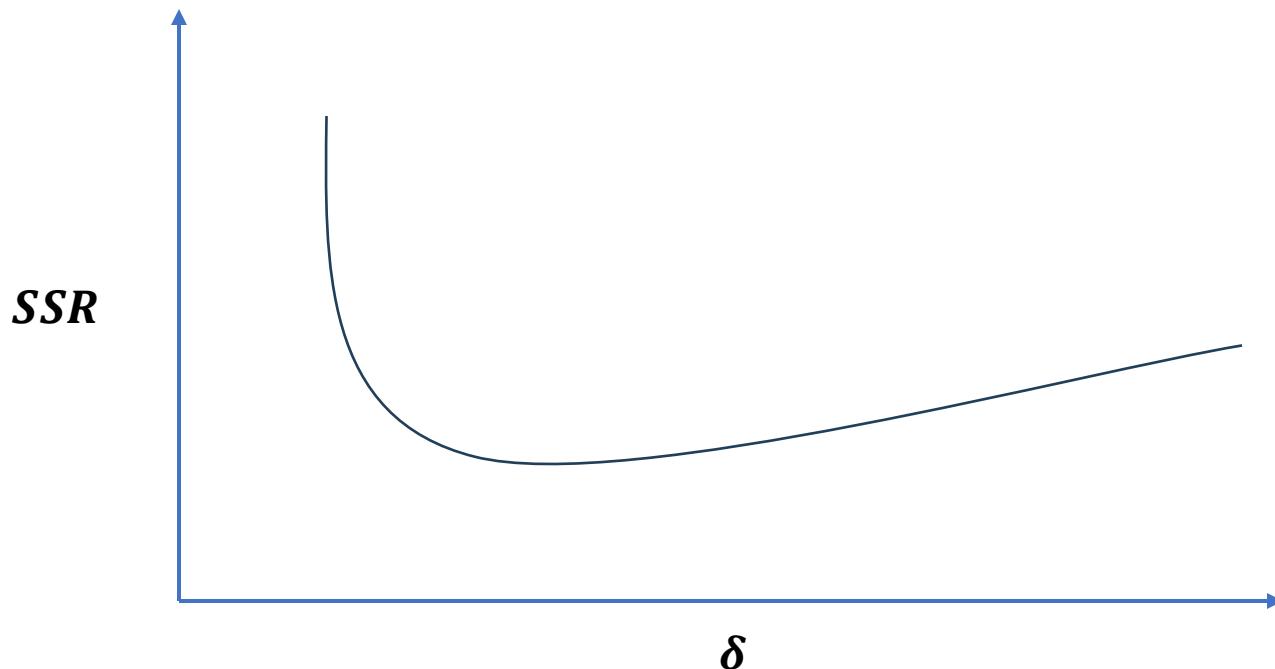
$$\text{SSR} = 1/2 \{(88 - \delta)^2 + (76 - \delta)^2 + (56 - \delta)^2\}$$

$$\frac{d\text{SSR}}{d\delta} = -(88 - \delta) - (76 - \delta) - (56 - \delta) = 0$$

$$\delta = \frac{88+76+56}{3} = 73.3 = \text{average of all sample' weights}$$

Gradient Boost: Behind The Scenes

- Step 1:
- Initialize a model with a constant value:
 - $F_0(x) = \operatorname{argmin}_{\delta} \sum_{i=1}^n L(y_i, \delta) = 73.3$



Gradient Boost: Behind The Scenes

- M is the number of trees
- n is the number of samples

Step 2: for $m = 1$ to M :

(A) Compute $r_{im} = - \left[\frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} \right]_{F(x)=F_{m-1}(x)}$ for $i = 1, \dots, n$

(B) Fit a regression tree to the r_{im} values and create terminal regions R_{jm} , for $j = 1 \dots J_m$

(C) For $j = 1 \dots J_m$ compute $\gamma_{jm} = \operatorname{argmin}_{\gamma} \sum_{x_i \in R_{ij}} L(y_i, F_{m-1}(x_i) + \gamma)$

(D) Update $F_m(x) = F_{m-1}(x) + \nu \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$

Gradient Boost: Behind The Scenes

(B) Fit a regression tree to the r_{im} values and create terminal regions R_{jm} , for $j = 1 \dots J_m$

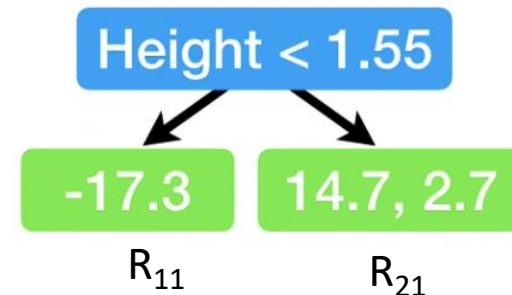
What?

How?

Why?

R is the residual error; m is the m-th tree; j is the j-th leaf

Height	Favorite Color	Gender	Weight	r_{i1}
1.6	Blue	Male	88	14.7
1.6	Green	Female	76	2.7
1.5	Blue	Female	56	-17.3



Step 1

$$F_0(x) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$$

(C) For $j = 1 \dots J_m$ compute $\gamma_{jm} = \operatorname{argmin}_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, F_{m-1}(x_i) + \gamma)$

$$F_{m-1}(x) = F_0(x) = 73.3$$

R is the residual error; m is the m-th tree; j is the i-th leaf

Height	Favorite Color	Gender	Weight	r_{i1}
1.6	Blue	Male	88	14.7
1.6	Green	Female	76	2.7
1.5	Blue	Female	56	-17.3

$$\gamma_{11} = \operatorname{argmin}_{\gamma} \left[\frac{1}{2} \{y_3 - (F_0(x_3) + \gamma)\}^2 \right]$$

$$\gamma_{11} = \operatorname{argmin}_{\gamma} \left[\frac{1}{2} \{56 - (73.3 + \gamma)\}^2 \right]$$

$$\gamma_{11} = \operatorname{argmin}_{\gamma} \left[\frac{1}{2} \{-17.3 - \gamma\}^2 \right]$$

$$\gamma_{11} = -17.3$$

$$\gamma_{21} = 8.7$$

Height < 1.55

-17.3 14.7, 2.7

$x_i \in R_{11}$

$x_i \in R_{21}$

Step 1

$$F_0(x) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$$

(C) For $j = 1 \dots J_m$ compute $\gamma_{jm} = \operatorname{argmin}_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, F_{m-1}(x_i) + \gamma)$

$$F_{m-1}(x) = F_0(x) = 73.3$$

R is the residual error; m is the m-th tree; j is the i-th leaf

Height	Favorite Color	Gender	Weight	r_{i1}
1.6	Blue	Male	88	14.7
1.6	Green	Female	76	2.7
1.5	Blue	Female	56	-17.3

$$\gamma_{21} = \operatorname{argmin}_{\gamma} \left[\frac{1}{2} \{y_1 - (F_0(x_1) + \gamma)\}^2 + \frac{1}{2} \{y_2 - (F_0(x_2) + \gamma)\}^2 \right]$$

$$\gamma_{11} = \operatorname{argmin}_{\gamma} \left[\frac{1}{2} \{88 - (73.3 + \gamma)\}^2 + \frac{1}{2} \{76 - (73.33 + \gamma)\}^2 \right]$$

$$\gamma_{11} = \operatorname{argmin}_{\gamma} \left[\frac{1}{2} \{14.7 - \gamma\}^2 + \frac{1}{2} \{2.7 - \gamma\}^2 \right]$$

$$\gamma_{21} = 8.7$$

Height < 1.55

-17.3, 14.7, 2.7

$x_i \in R_{11}$ $x_i \in R_{21}$

$$\gamma_{11} = -17.3$$

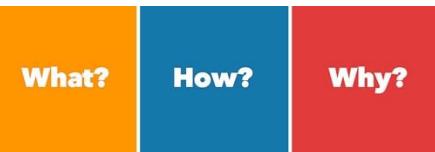
$$\gamma_{21} = 8.7$$

Giá trị trung
bình của 2
samples

Step 2

Step 1

$$x_i \in R_{11}$$



$$F_{m-1}(x) = F_0(x) = 73.3$$

Update $F_m(x) = F_{m-1}(x) + \nu \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$

R is the residual error; m is the m-th tree; j is the i-th leaf

Height	Favorite Color	Gender	Weight	r_{i1}	$F_1(x)$
1.6	Blue	Male	88	14.7	74.2
1.6	Green	Female	76	2.7	74.2
1.5	Blue	Female	56	-17.3	71.6

Height < 1.55

-17.3 14.7, 2.7

$$x_i \in R_{11}$$

$$x_i \in R_{21}$$

$$\gamma_{11} = -17.3$$

$$\gamma_{21} = 8.7$$

Repeat for the next $m + 1$ iteration

Step 2: for $m = 1$ to M :

(A) Compute $r_{im} = - \left[\frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} \right]_{F(x)=F_{m-1}(x)}$ for $i = 1, \dots, n$

(B) Fit a regression tree to the r_{im} values and create terminal regions R_{jm} , for $j = 1 \dots J_m$

(C) For $j = 1 \dots J_m$ compute $\gamma_{jm} = \operatorname{argmin}_{\gamma} \sum_{x_i \in R_{ij}} L(y_i, F_{m-1}(x_i) + \gamma)$

(D) Update $F_m(x) = F_{m-1}(x) + \nu \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$

Summary

Input: training set $\{(x_i, y_i)\}_{i=1}^n$, a differentiable loss function $L(y, F(x))$, number of iterations M .

Algorithm:

1. Initialize model with a constant value:

$$F_0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, \gamma).$$

2. For $m = 1$ to M :

1. Compute so-called *pseudo-residuals*:

$$r_{im} = - \left[\frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} \right]_{F(x)=F_{m-1}(x)} \quad \text{for } i = 1, \dots, n.$$

2. Fit a base learner (e.g. tree) $h_m(x)$ to pseudo-residuals, i.e. train it using the training set $\{(x_i, r_{im})\}_{i=1}^n$.

3. Compute multiplier γ_m by solving the following *one-dimensional optimization* problem:

$$\gamma_m = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, F_{m-1}(x_i) + \gamma h_m(x_i)).$$

4. Update the model:

$$F_m(x) = F_{m-1}(x) + \gamma_m h_m(x).$$

3. Output $F_M(x)$.

Outline

- Boosting Techniques
- AdaBoost Clearly Explain
- Gradient Boost Clearly Explain
- Time Series Data: Predicting Energy Consumption
- Summary



Time Series Forecasting

We will focus on the energy consumption problem, where given a sufficiently large dataset of the daily energy consumption of different households in a city, we are tasked to predict as accurately as possible the future energy demands.

london_energy

LCLid	Date	KWH
MAC000002	2012-10-12	7.098
MAC000002	2012-10-13	11.087
MAC000002	2012-10-14	13.223
MAC000002	2012-10-15	10.257
MAC000002	2012-10-16	9.769
MAC000002	2012-10-17	10.885
MAC000002	2012-10-18	10.751
MAC000002	2012-10-19	8.431
MAC000002	2012-10-20	17.578
MAC000002	2012-10-21	24.49
MAC000002	2012-10-22	18.885
MAC000002	2012-10-23	10.485
MAC000002	2012-10-24	15.537

Preprocessing

```
▶ import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv("/content/drive/MyDrive/AI02024/london_energy.csv")
print(df.isna().sum())
df.head()
```

```
→ LCLid      0
Date       0
KWH       0
dtype: int64
```

	LCLid	Date	KWH
0	MAC000002	2012-10-12	7.098
1	MAC000002	2012-10-13	11.087
2	MAC000002	2012-10-14	13.223
3	MAC000002	2012-10-15	10.257
4	MAC000002	2012-10-16	9.769

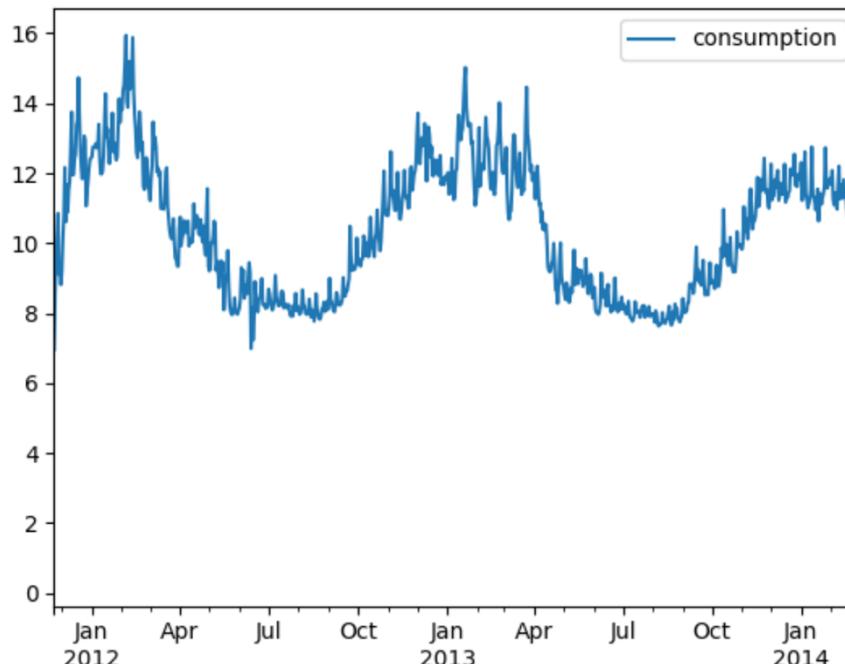
Time Series Forecasting

We will focus on the energy consumption problem, where given a sufficiently large dataset of the daily energy consumption of different households in a city, we are tasked to predict as accurately as possible the future energy demands.

Consumption changes through the years

```
▶ df_avg_consumption.plot(x="date", y="consumption")
```

```
↳ <Axes: xlabel='date'>
```



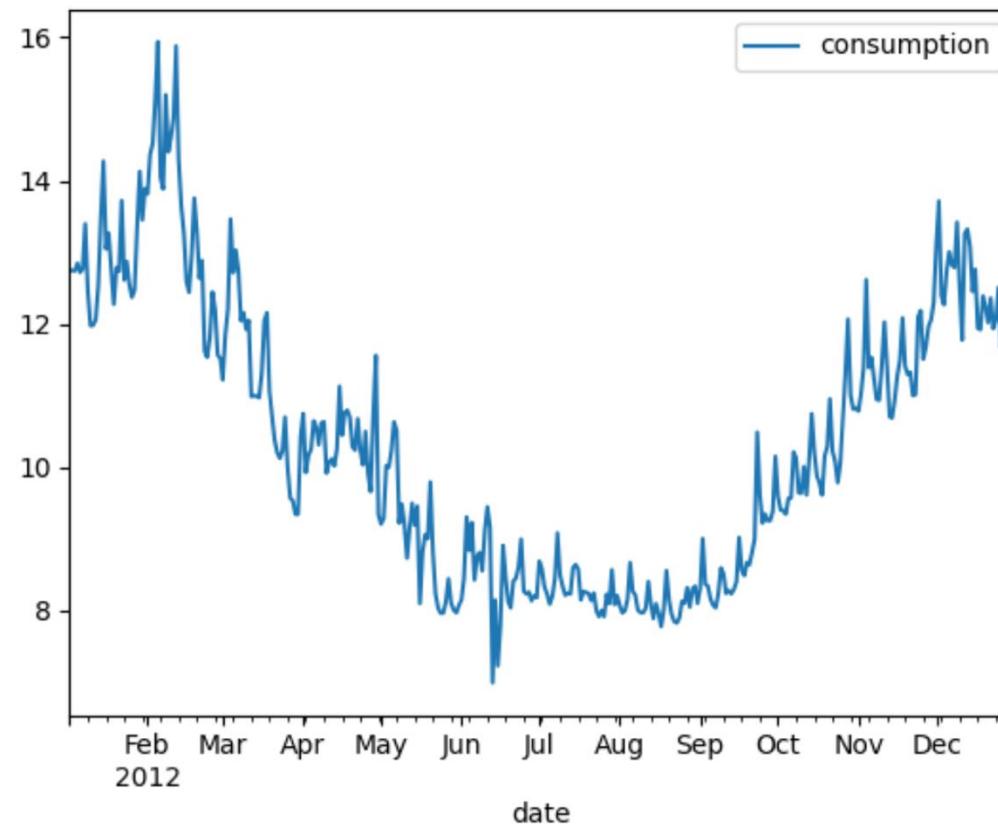
During the winter months we observe high demands in energy, while throughout the summer the consumption is at the lowest levels.

Time Series Forecasting

Visualize the fluctuation in the span of a year we can do

```
▶ df_avg_consumption.query("date > '2012-01-01' & date < '2013-01-01'").plot(x="date", y="consumption")
```

```
⤵ <Axes: xlabel='date'>
```



Time Series Forecasting

We have only one feature: the full date. We can extract different features based on the full date such as the day of the week, the day of the year, the month and others

```
▶ df_avg_consumption["day_of_week"] = df_avg_consumption["date"].dt.dayofweek
df_avg_consumption["day_of_year"] = df_avg_consumption["date"].dt.dayofyear
df_avg_consumption["month"] = df_avg_consumption["date"].dt.month
df_avg_consumption["quarter"] = df_avg_consumption["date"].dt.quarter
df_avg_consumption["year"] = df_avg_consumption["date"].dt.year

df_avg_consumption.head()
```

	date	consumption	day_of_week	day_of_year	month	quarter	year	grid icon	bar chart icon
0	2011-11-23	6.952692	2	327	11	4	2011		
1	2011-11-24	8.536480	3	328	11	4	2011		
2	2011-11-25	9.499781	4	329	11	4	2011		
3	2011-11-26	10.267707	5	330	11	4	2011		
4	2011-11-27	10.850805	6	331	11	4	2011		

Time Series Forecasting

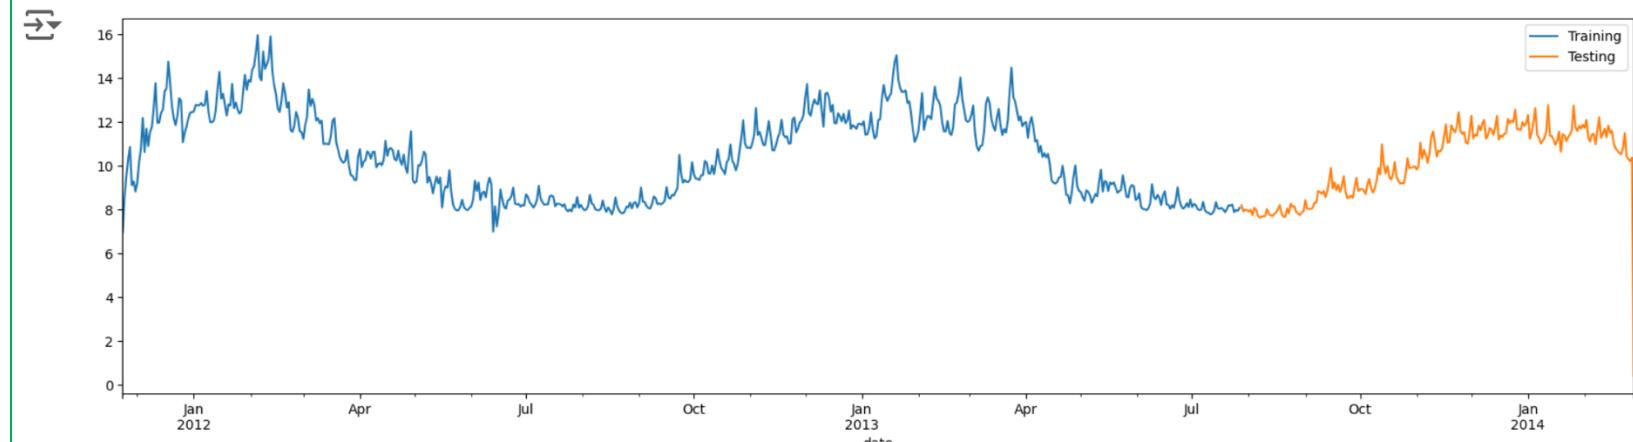
The dataset contains almost 2.5 years of data, so for the testing set we will use only the last 6 months

```
[6] training_mask = df_avg_consumption["date"] < "2013-07-28"
    training_data = df_avg_consumption.loc[training_mask]
    print(training_data.shape)

    testing_mask = df_avg_consumption["date"] >= "2013-07-28"
    testing_data = df_avg_consumption.loc[testing_mask]
    print(testing_data.shape)
```

```
→ (613, 7)
(216, 7)
```

```
▶ figure, ax = plt.subplots(figsize=(20, 5))
    training_data.plot(ax=ax, label="Training", x="date", y="consumption")
    testing_data.plot(ax=ax, label="Testing", x="date", y="consumption")
    plt.show()
```



Time Series Forecasting

Prepare dataset for Training and Testing

```
[8] # Dropping unnecessary `date` column
training_data = training_data.drop(columns=["date"])
testing_dates = testing_data["date"]
testing_data = testing_data.drop(columns=["date"])

X_train = training_data[["day_of_week", "day_of_year", "month", "quarter", "year"]]
y_train = training_data["consumption"]

X_test = testing_data[["day_of_week", "day_of_year", "month", "quarter", "year"]]
y_test = testing_data["consumption"]
```

```
[12] from sklearn.metrics import mean_absolute_error, mean_absolute_percentage_error,\n    mean_squared_error

def evaluate_model(y_test, prediction):
    print(f"MAE: {mean_absolute_error(y_test, prediction)}")
    print(f"MSE: {mean_squared_error(y_test, prediction)}")
    print(f"MAPE: {mean_absolute_percentage_error(y_test, prediction)}")

def plot_predictions(testing_dates, y_test, prediction):
    df_test = pd.DataFrame({"date": testing_dates, "actual": y_test, "prediction": prediction })
    figure, ax = plt.subplots(figsize=(10, 5))
    df_test.plot(ax=ax, label="Actual", x="date", y="actual")
    df_test.plot(ax=ax, label="Prediction", x="date", y="prediction")
    plt.legend(["Actual", "Prediction"])
    plt.show()
```

Performance Evaluation and Visualization

Time Series Forecasting

Gradient Boosting For Training

```
[14] from sklearn.model_selection import TimeSeriesSplit, GridSearchCV
     from sklearn.ensemble import GradientBoostingRegressor

     # XGBoost
     cv_split = TimeSeriesSplit(n_splits=4, test_size=100)
     model = GradientBoostingRegressor()
     parameters = {
         "max_features": [3, 4, 5],
         "learning_rate": [0.01, 0.05],
         "n_estimators": [100, 300]
     }

     grid_search = GridSearchCV(estimator=model, cv=cv_split, param_grid=parameters)
     grid_search.fit(X_train, y_train)
```

Time Series Forecasting

Gradient Boosting For Predicting

```
▶ # Evaluating GridSearch results
prediction = grid_search.predict(X_test)
plot_predictions(testing_dates, y_test, prediction)
evaluate_model(y_test, prediction)
```

